

# FORECASTING DIRECTIONAL MOVEMENTS USING LSTM

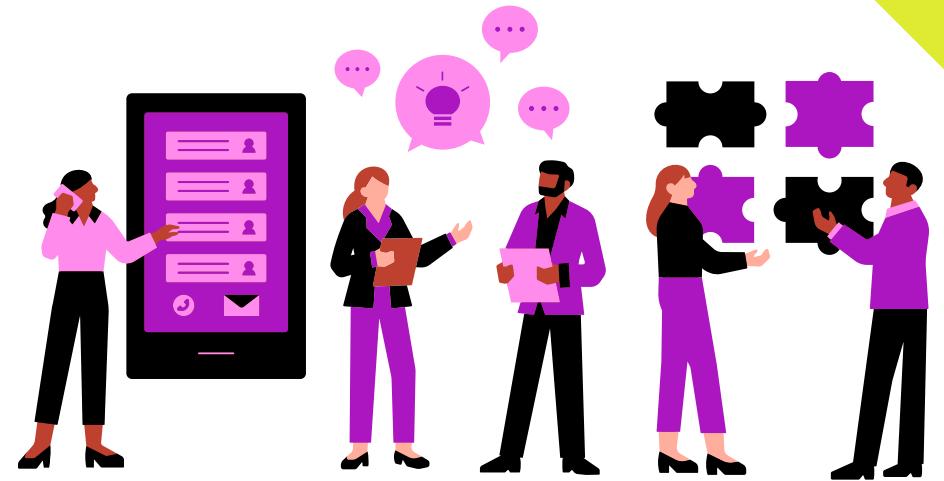
SOFTWARE ENGINEERING PROJECT

PRES E N T E D   B Y   G R O U P - 1 1





# MEET OUR TEAM



**B. Syam Sai Krishna-**  
**AM.EN.U4CSE21016**

Project Lead



**A.S.V Prasasya-**  
**AM.EN.U4CSE21067**

Business Analyst



**Purnima R -**  
**AM.EN.U4CSE21046**

Scrum Master





# INTRODUCTION

- The forex (foreign exchange) market involves the trading of currencies globally. It is the largest and most liquid financial market in the world, with a daily turnover of trillions of dollars. Having accurate forecasts can provide traders with a competitive advantage in the market. LSTM networks contain memory cells that can store information over extended periods, allowing them to retain important context.

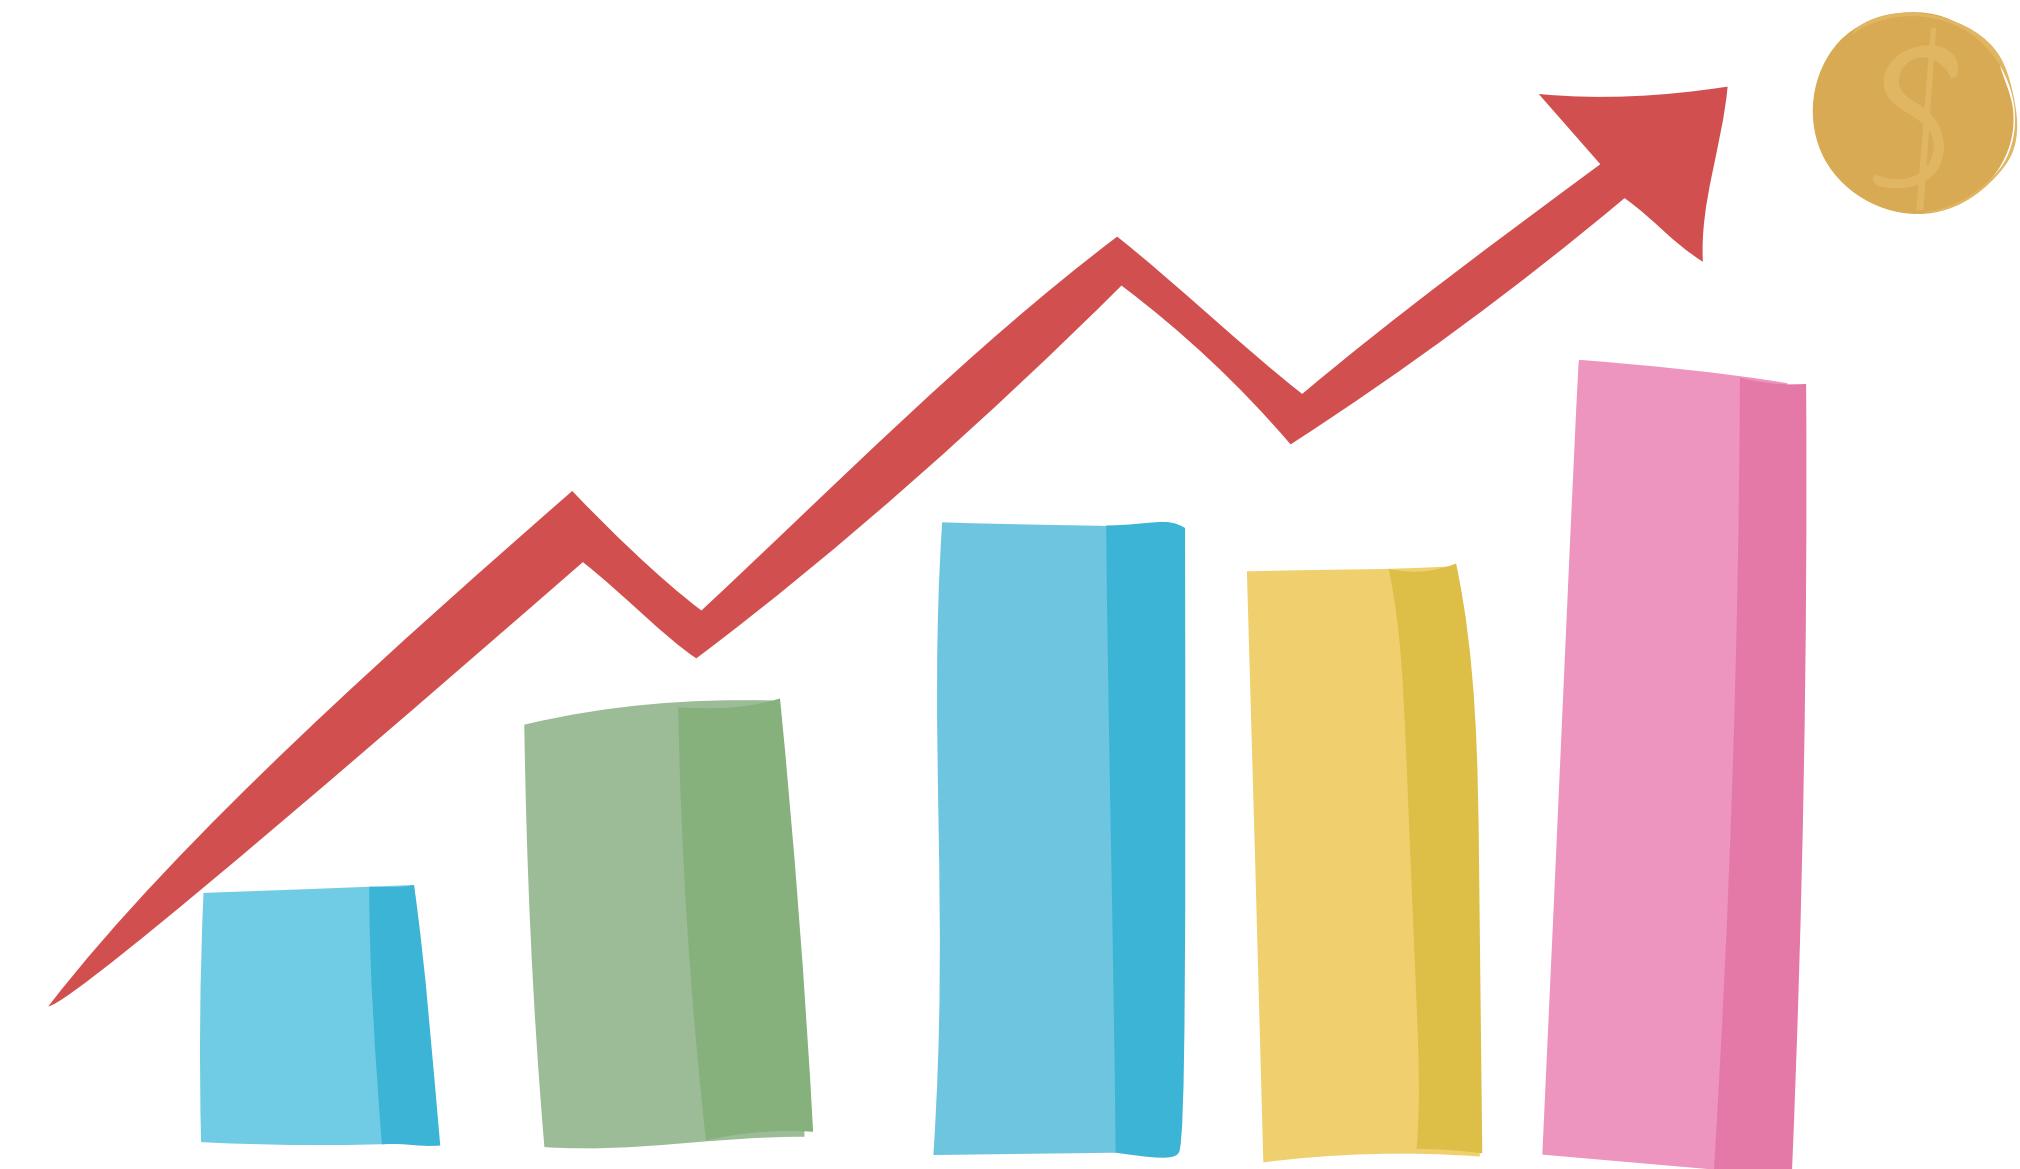
# PROBLEM STATEMENT

- Forecasting directional movement in Forex (foreign exchange) data is a challenging task due to the inherent complexity and volatility of the currency markets.
- Traders and investors need accurate predictions of whether a currency pair will move up or down in order to make informed decisions and maximize profits. Traditional time series forecasting techniques often struggle to capture the nonlinear and dynamic nature of Forex data, making it difficult to achieve reliable predictions.





- USE OF DIFFERENT TECHNICAL INDICATORS
- VISUALIZATION FOR ANALYSIS
- IMPORTANCE OF TECHNICAL INDICATORS IN RESEARCH



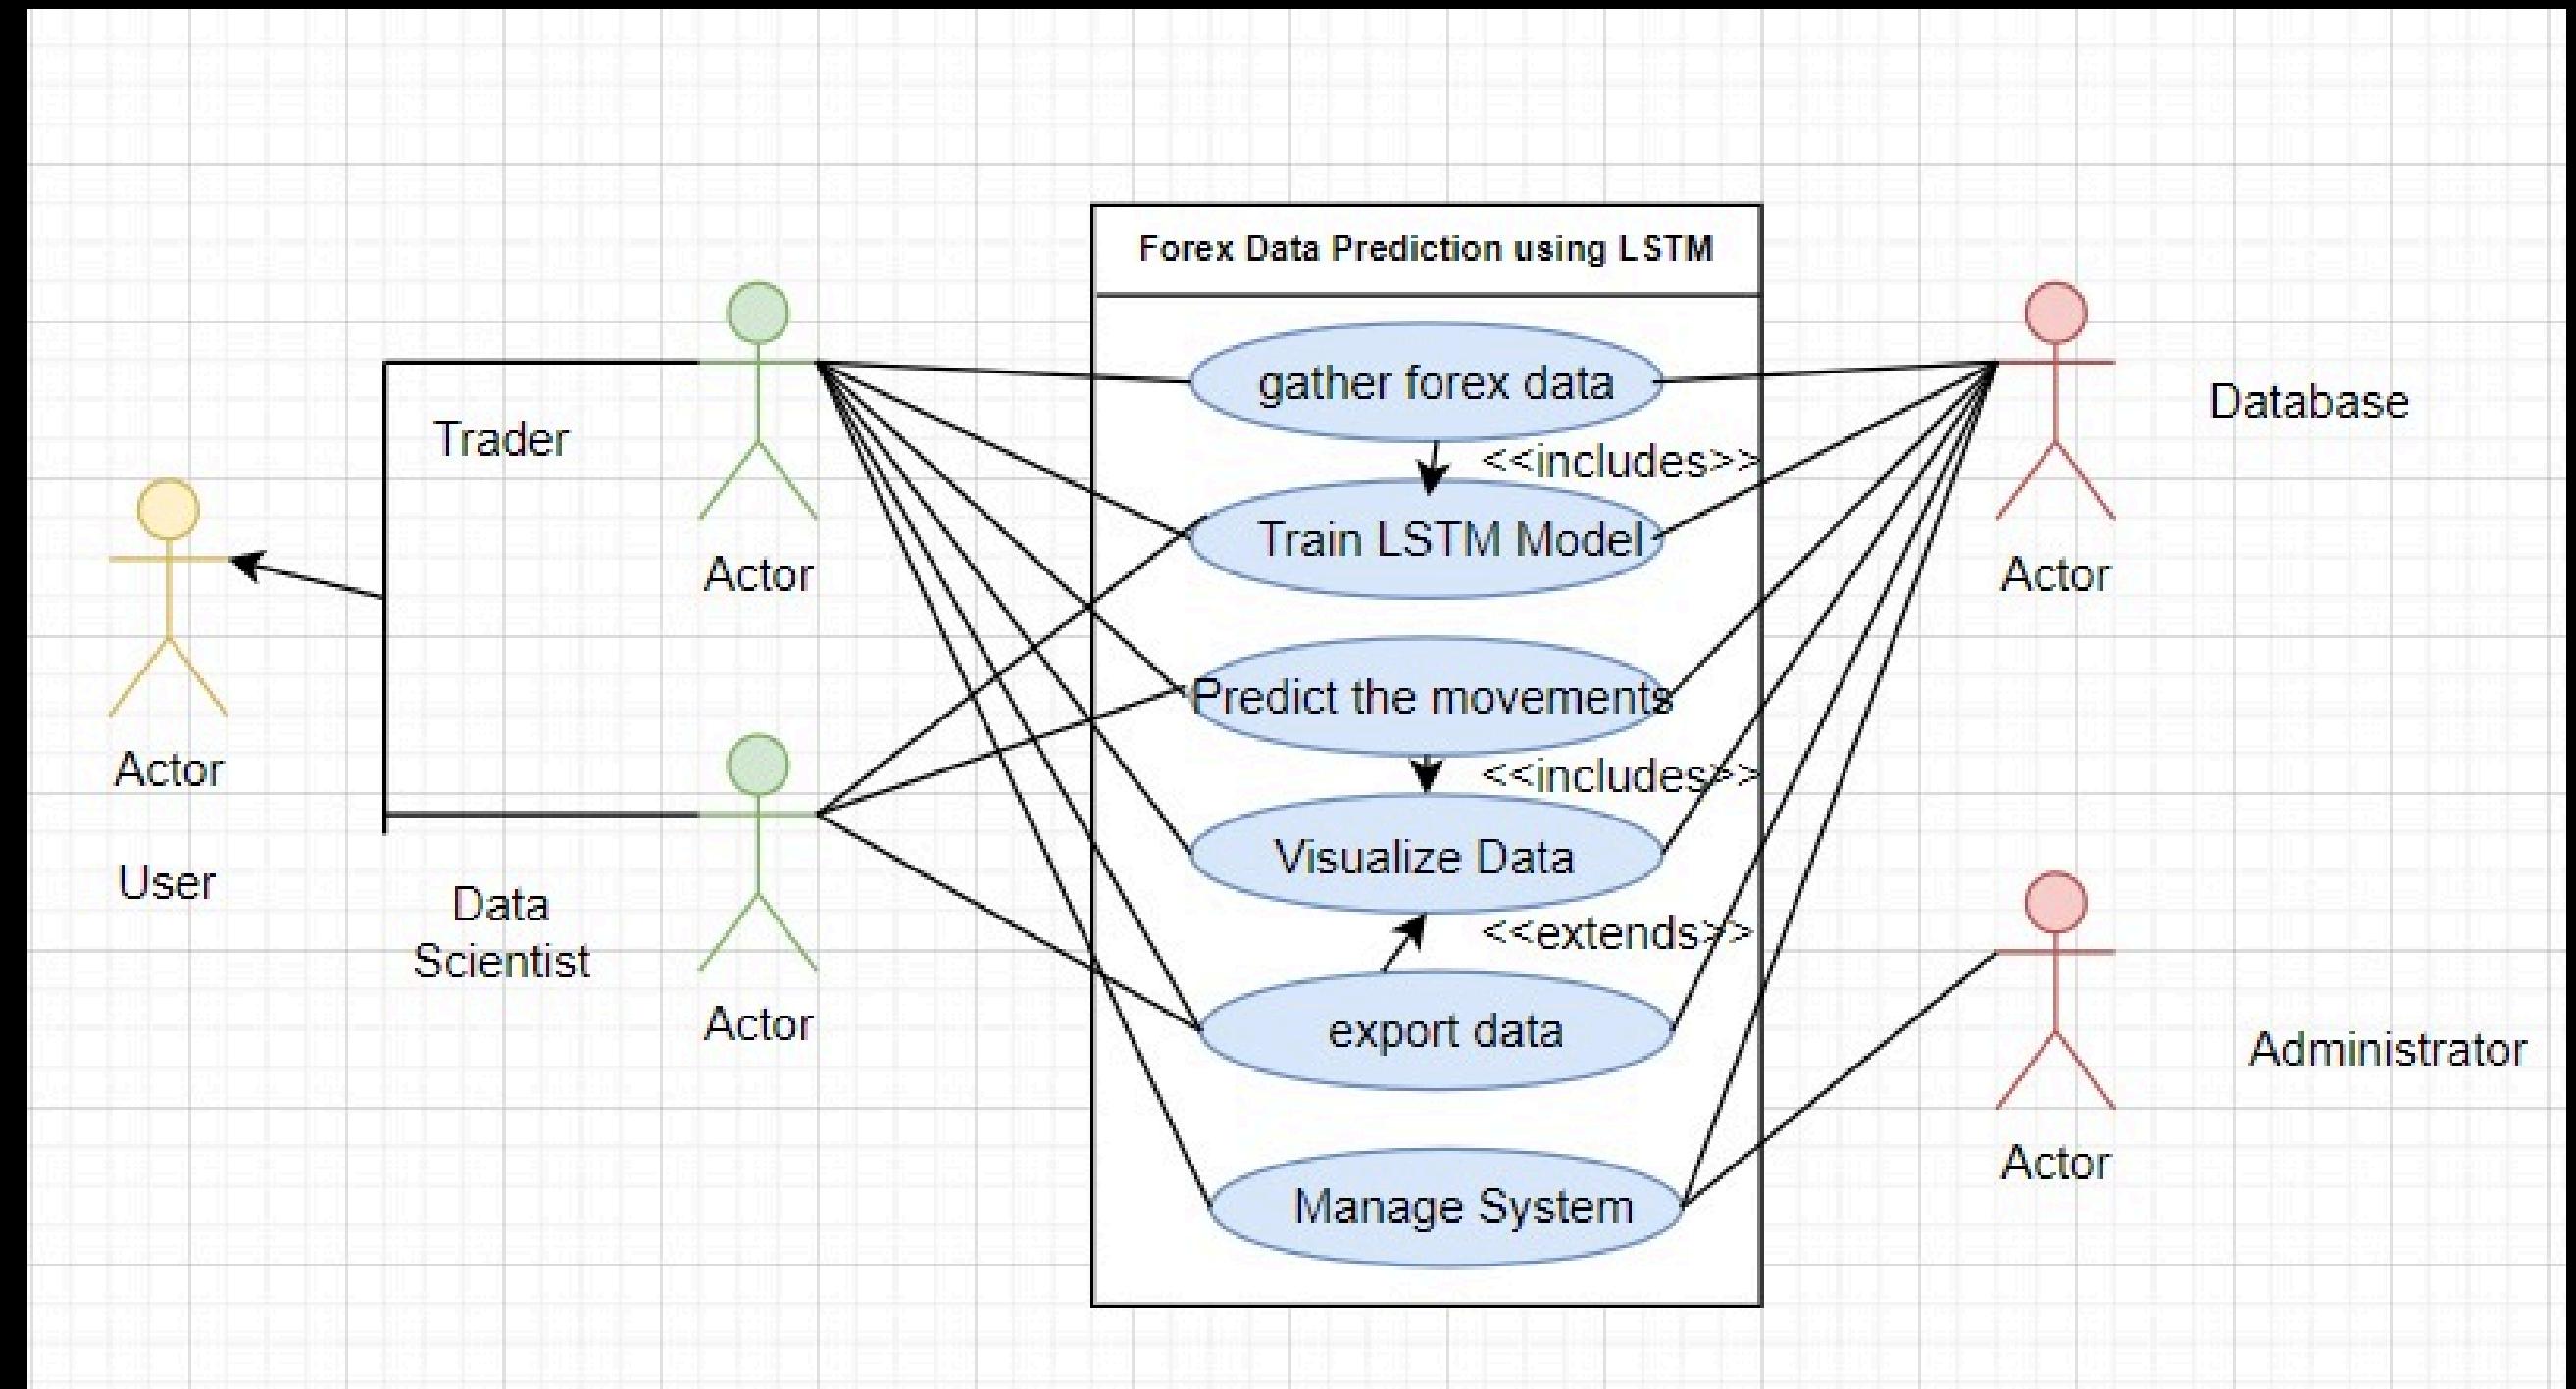
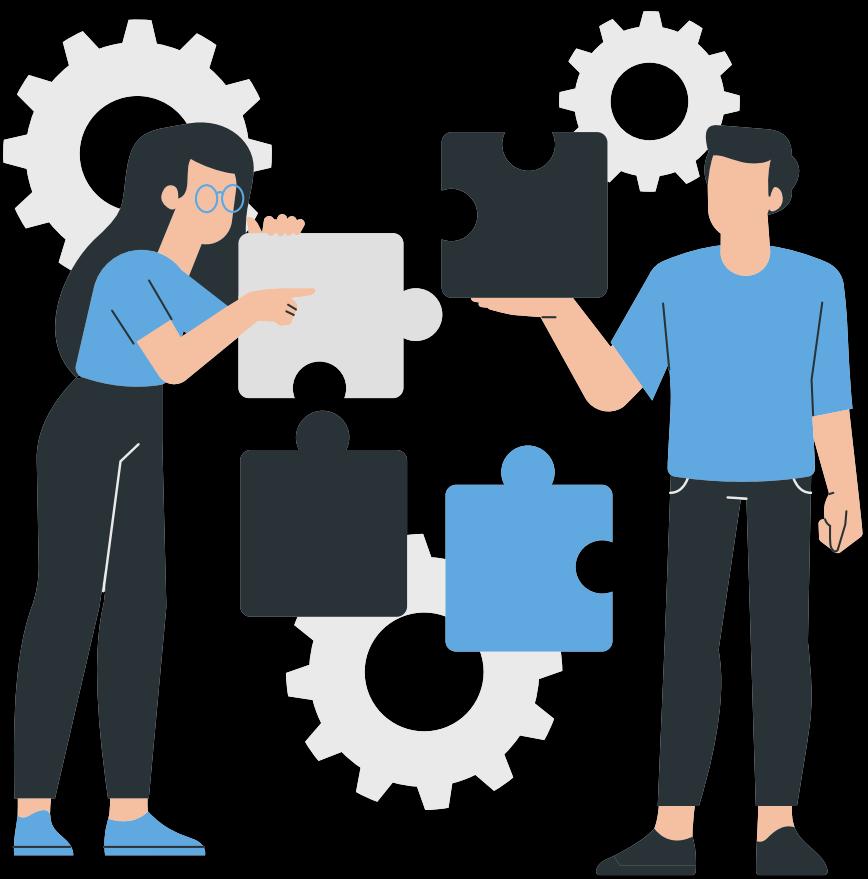
# MODEL- DESIGN PHASE

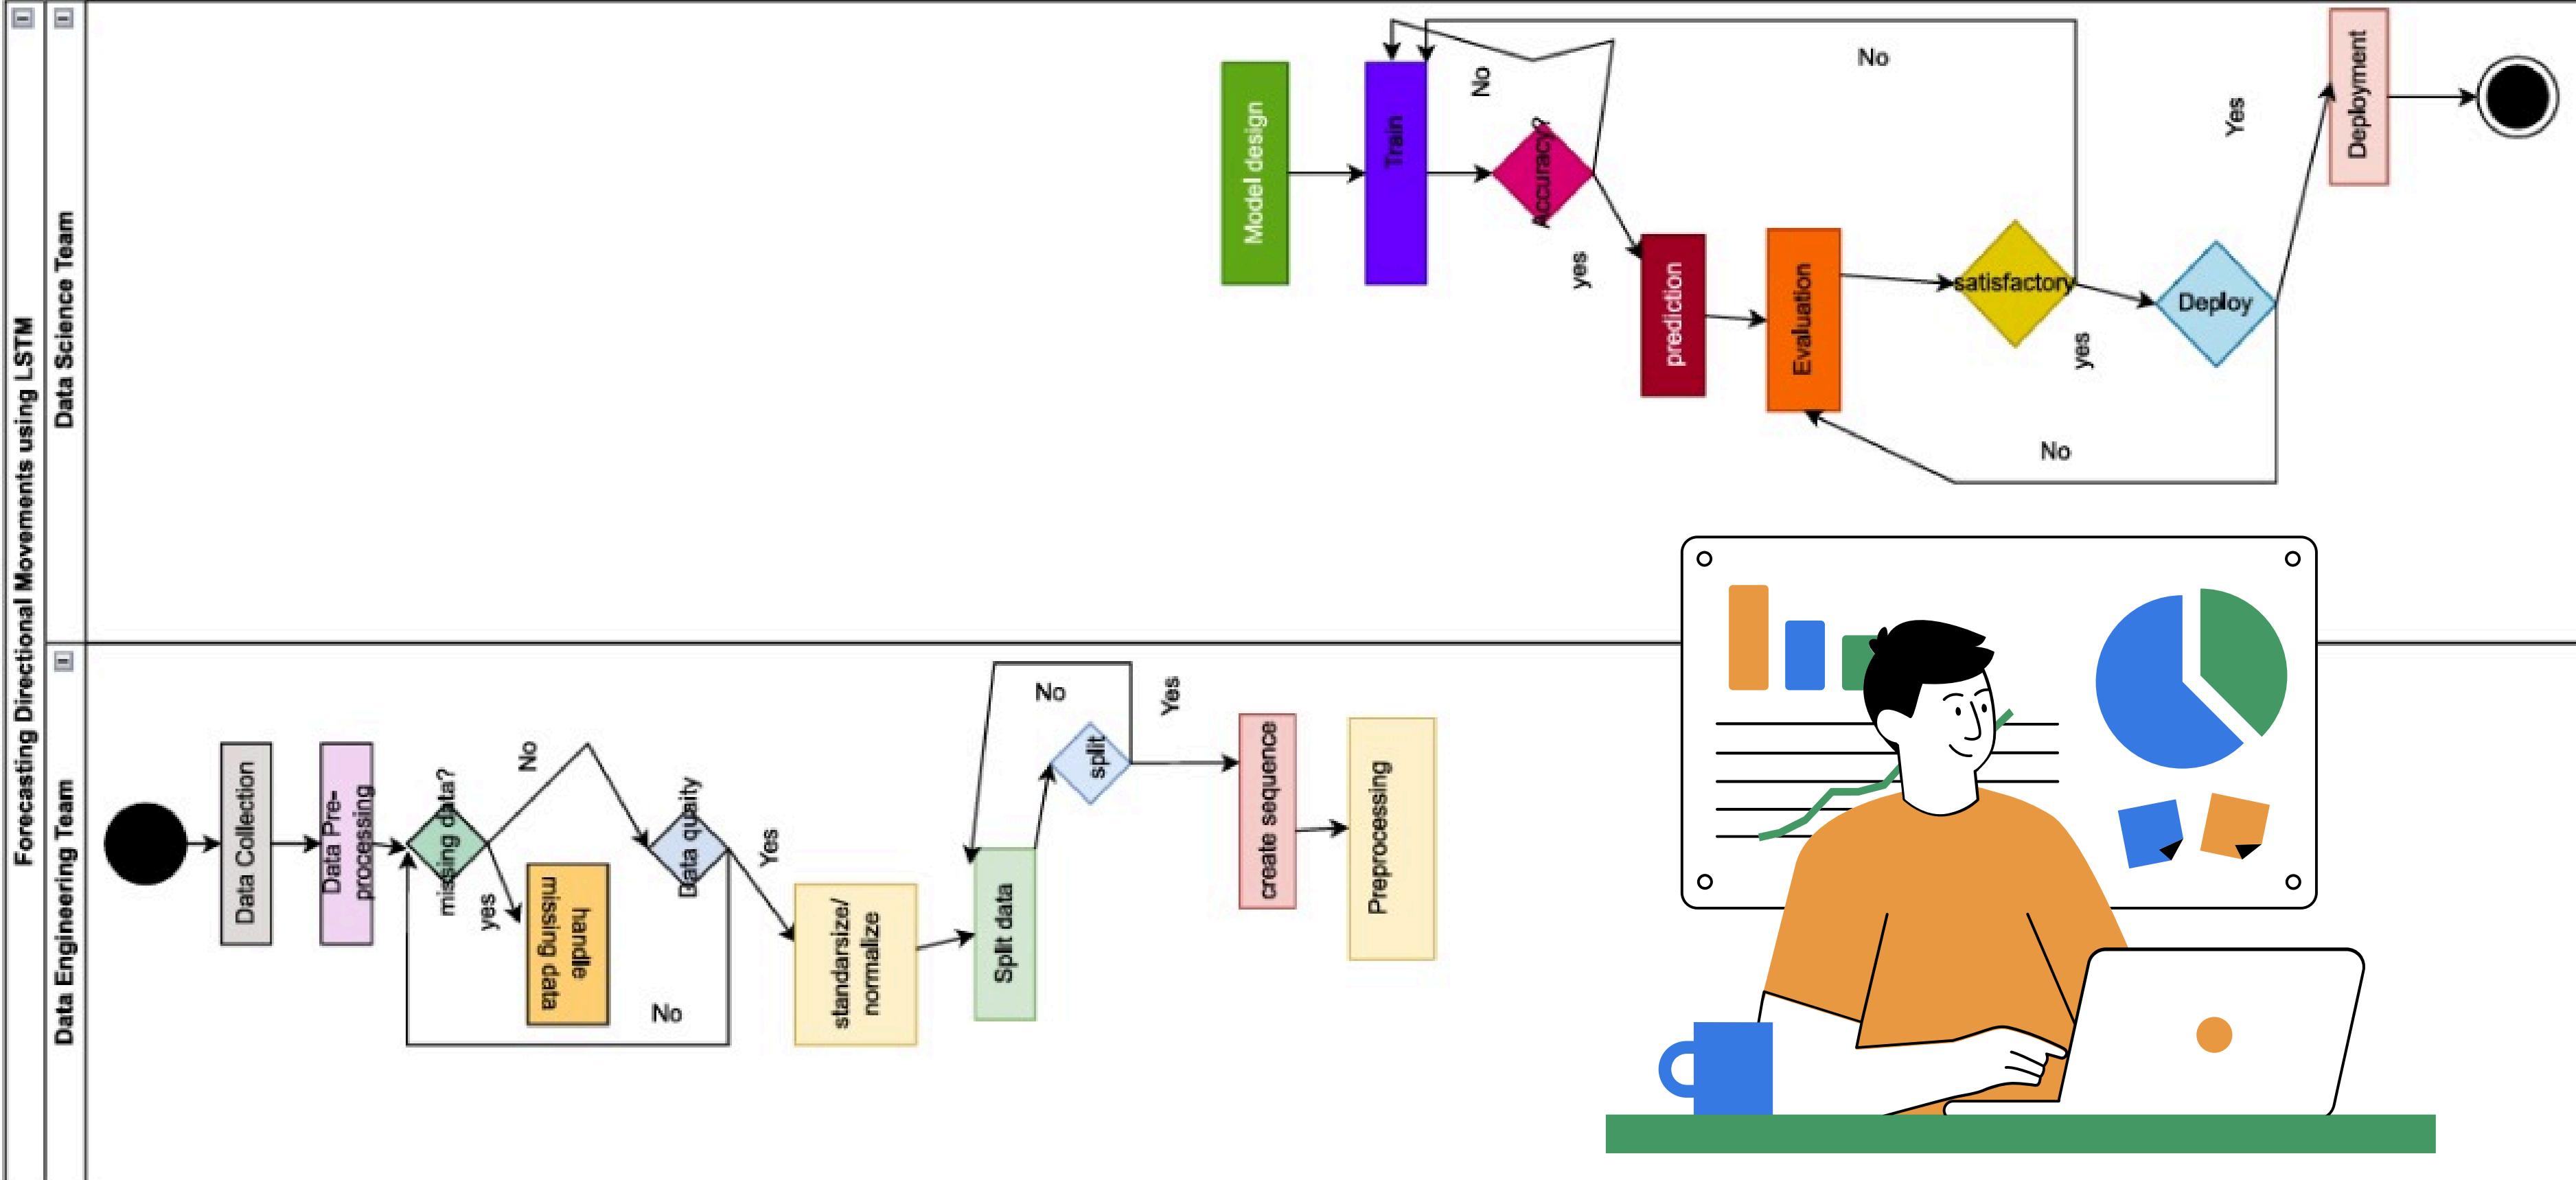


- 01** USE-CASE
- 02** ACTIVITY
- 03** CLASS
- 04** DATA FLOW
- 05** SEQUENCE



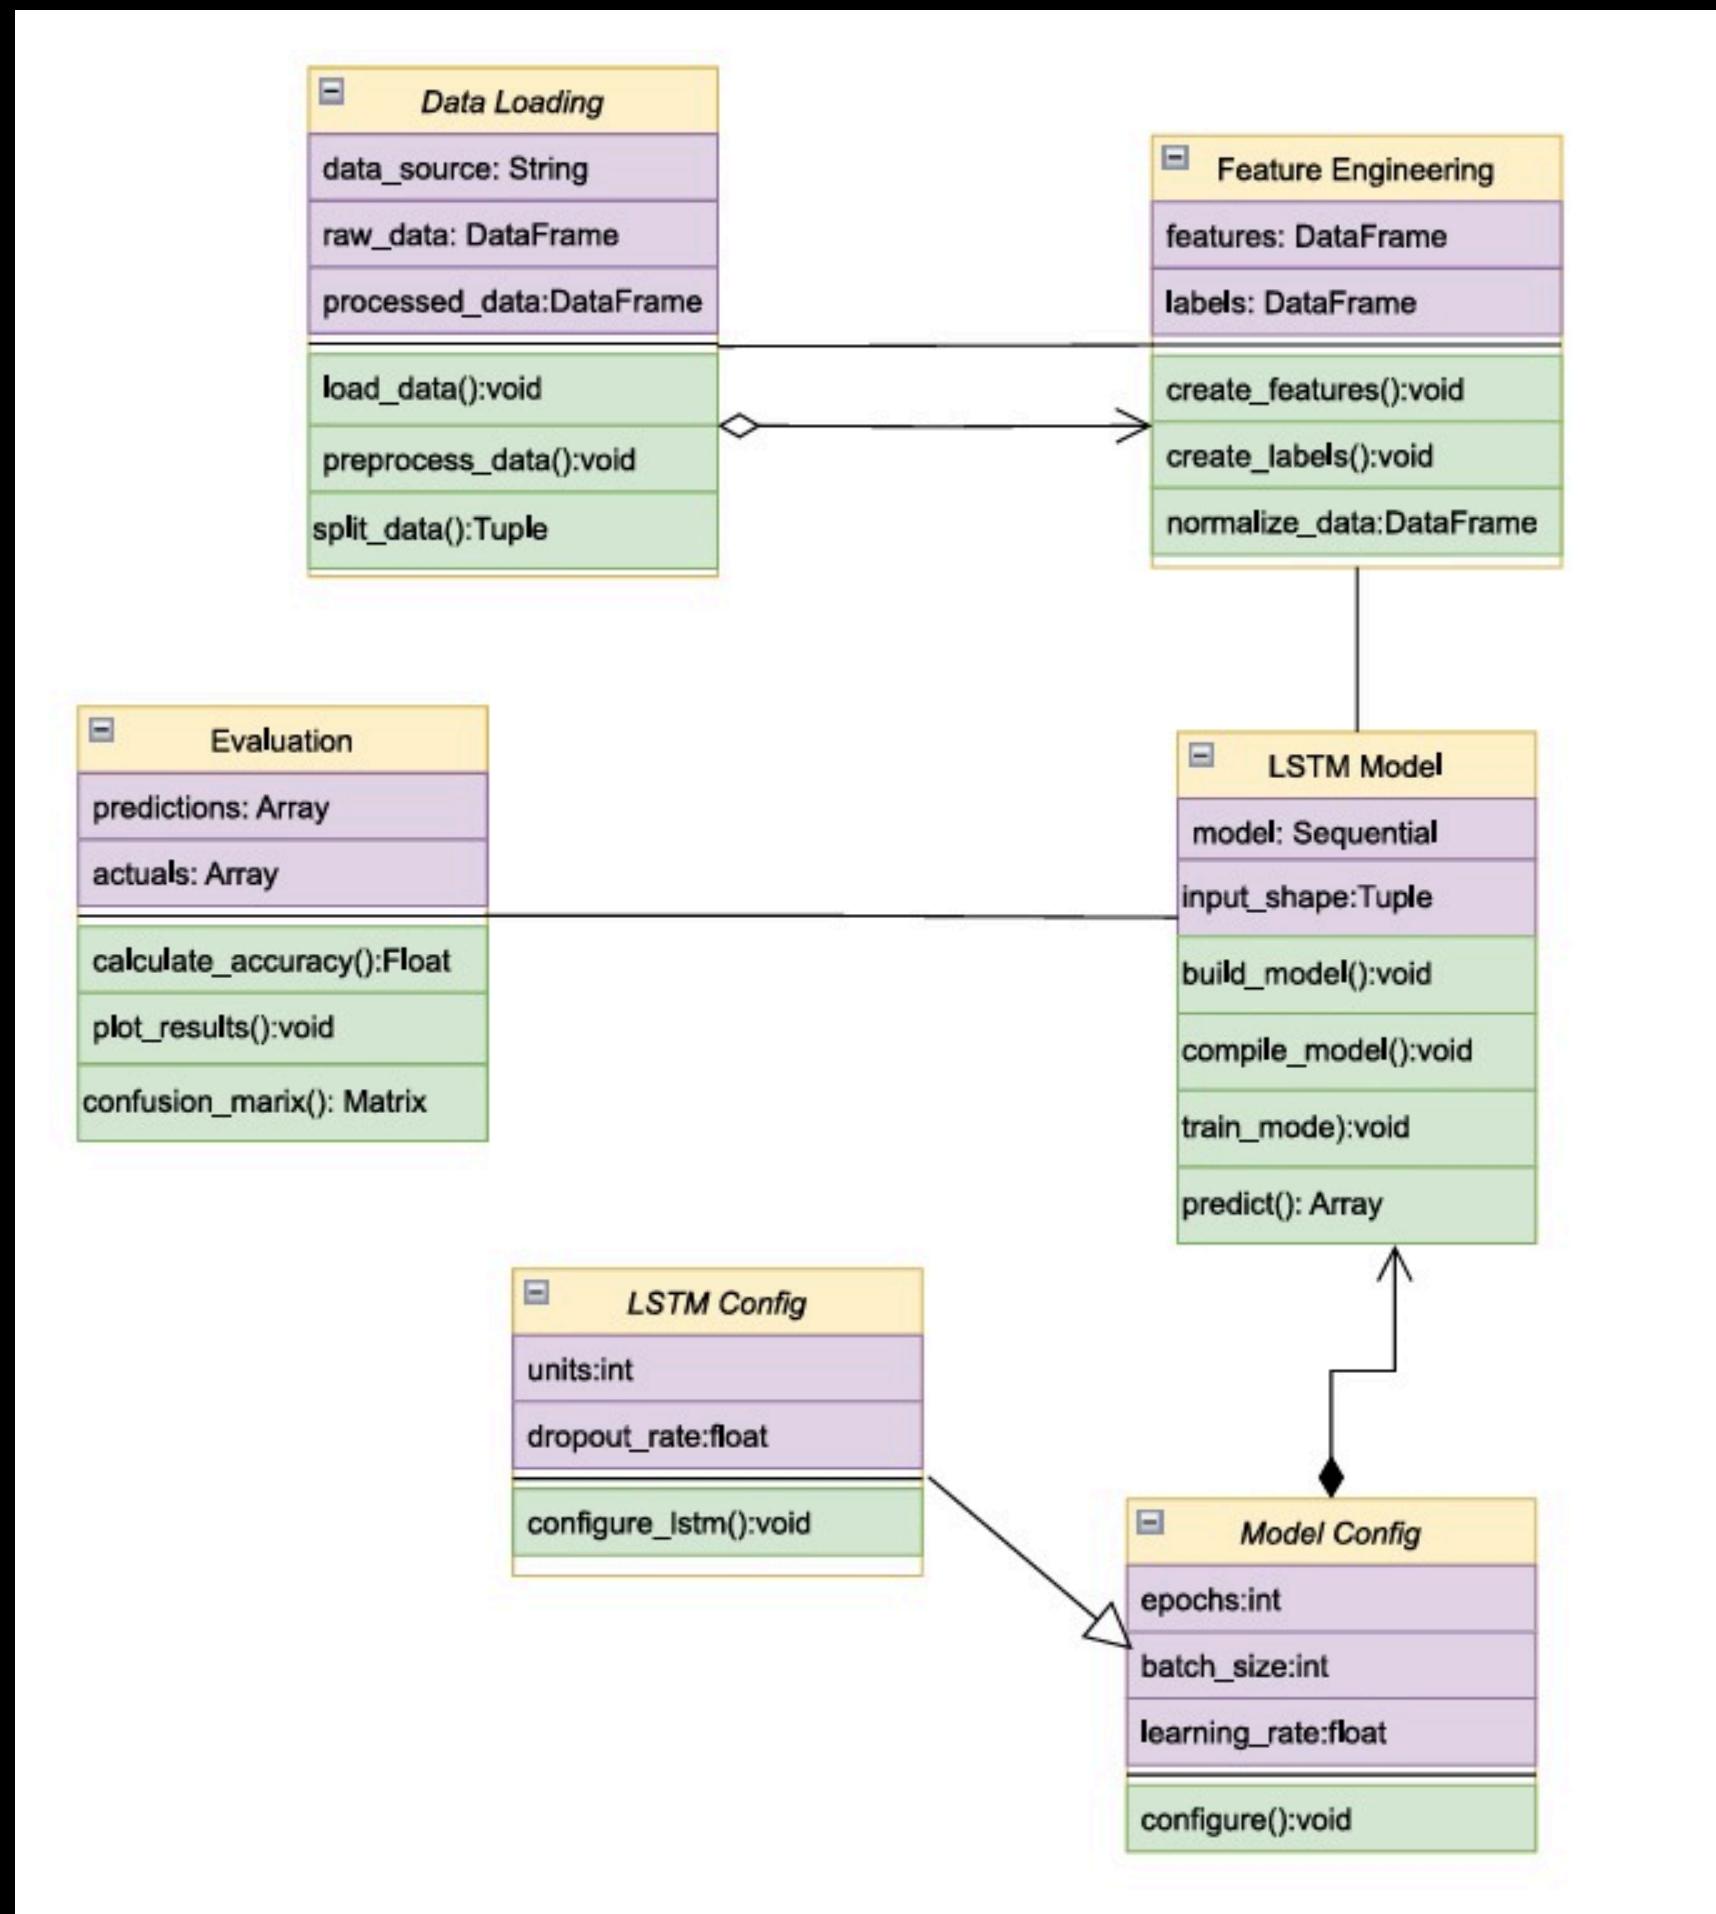
# USE-CASE



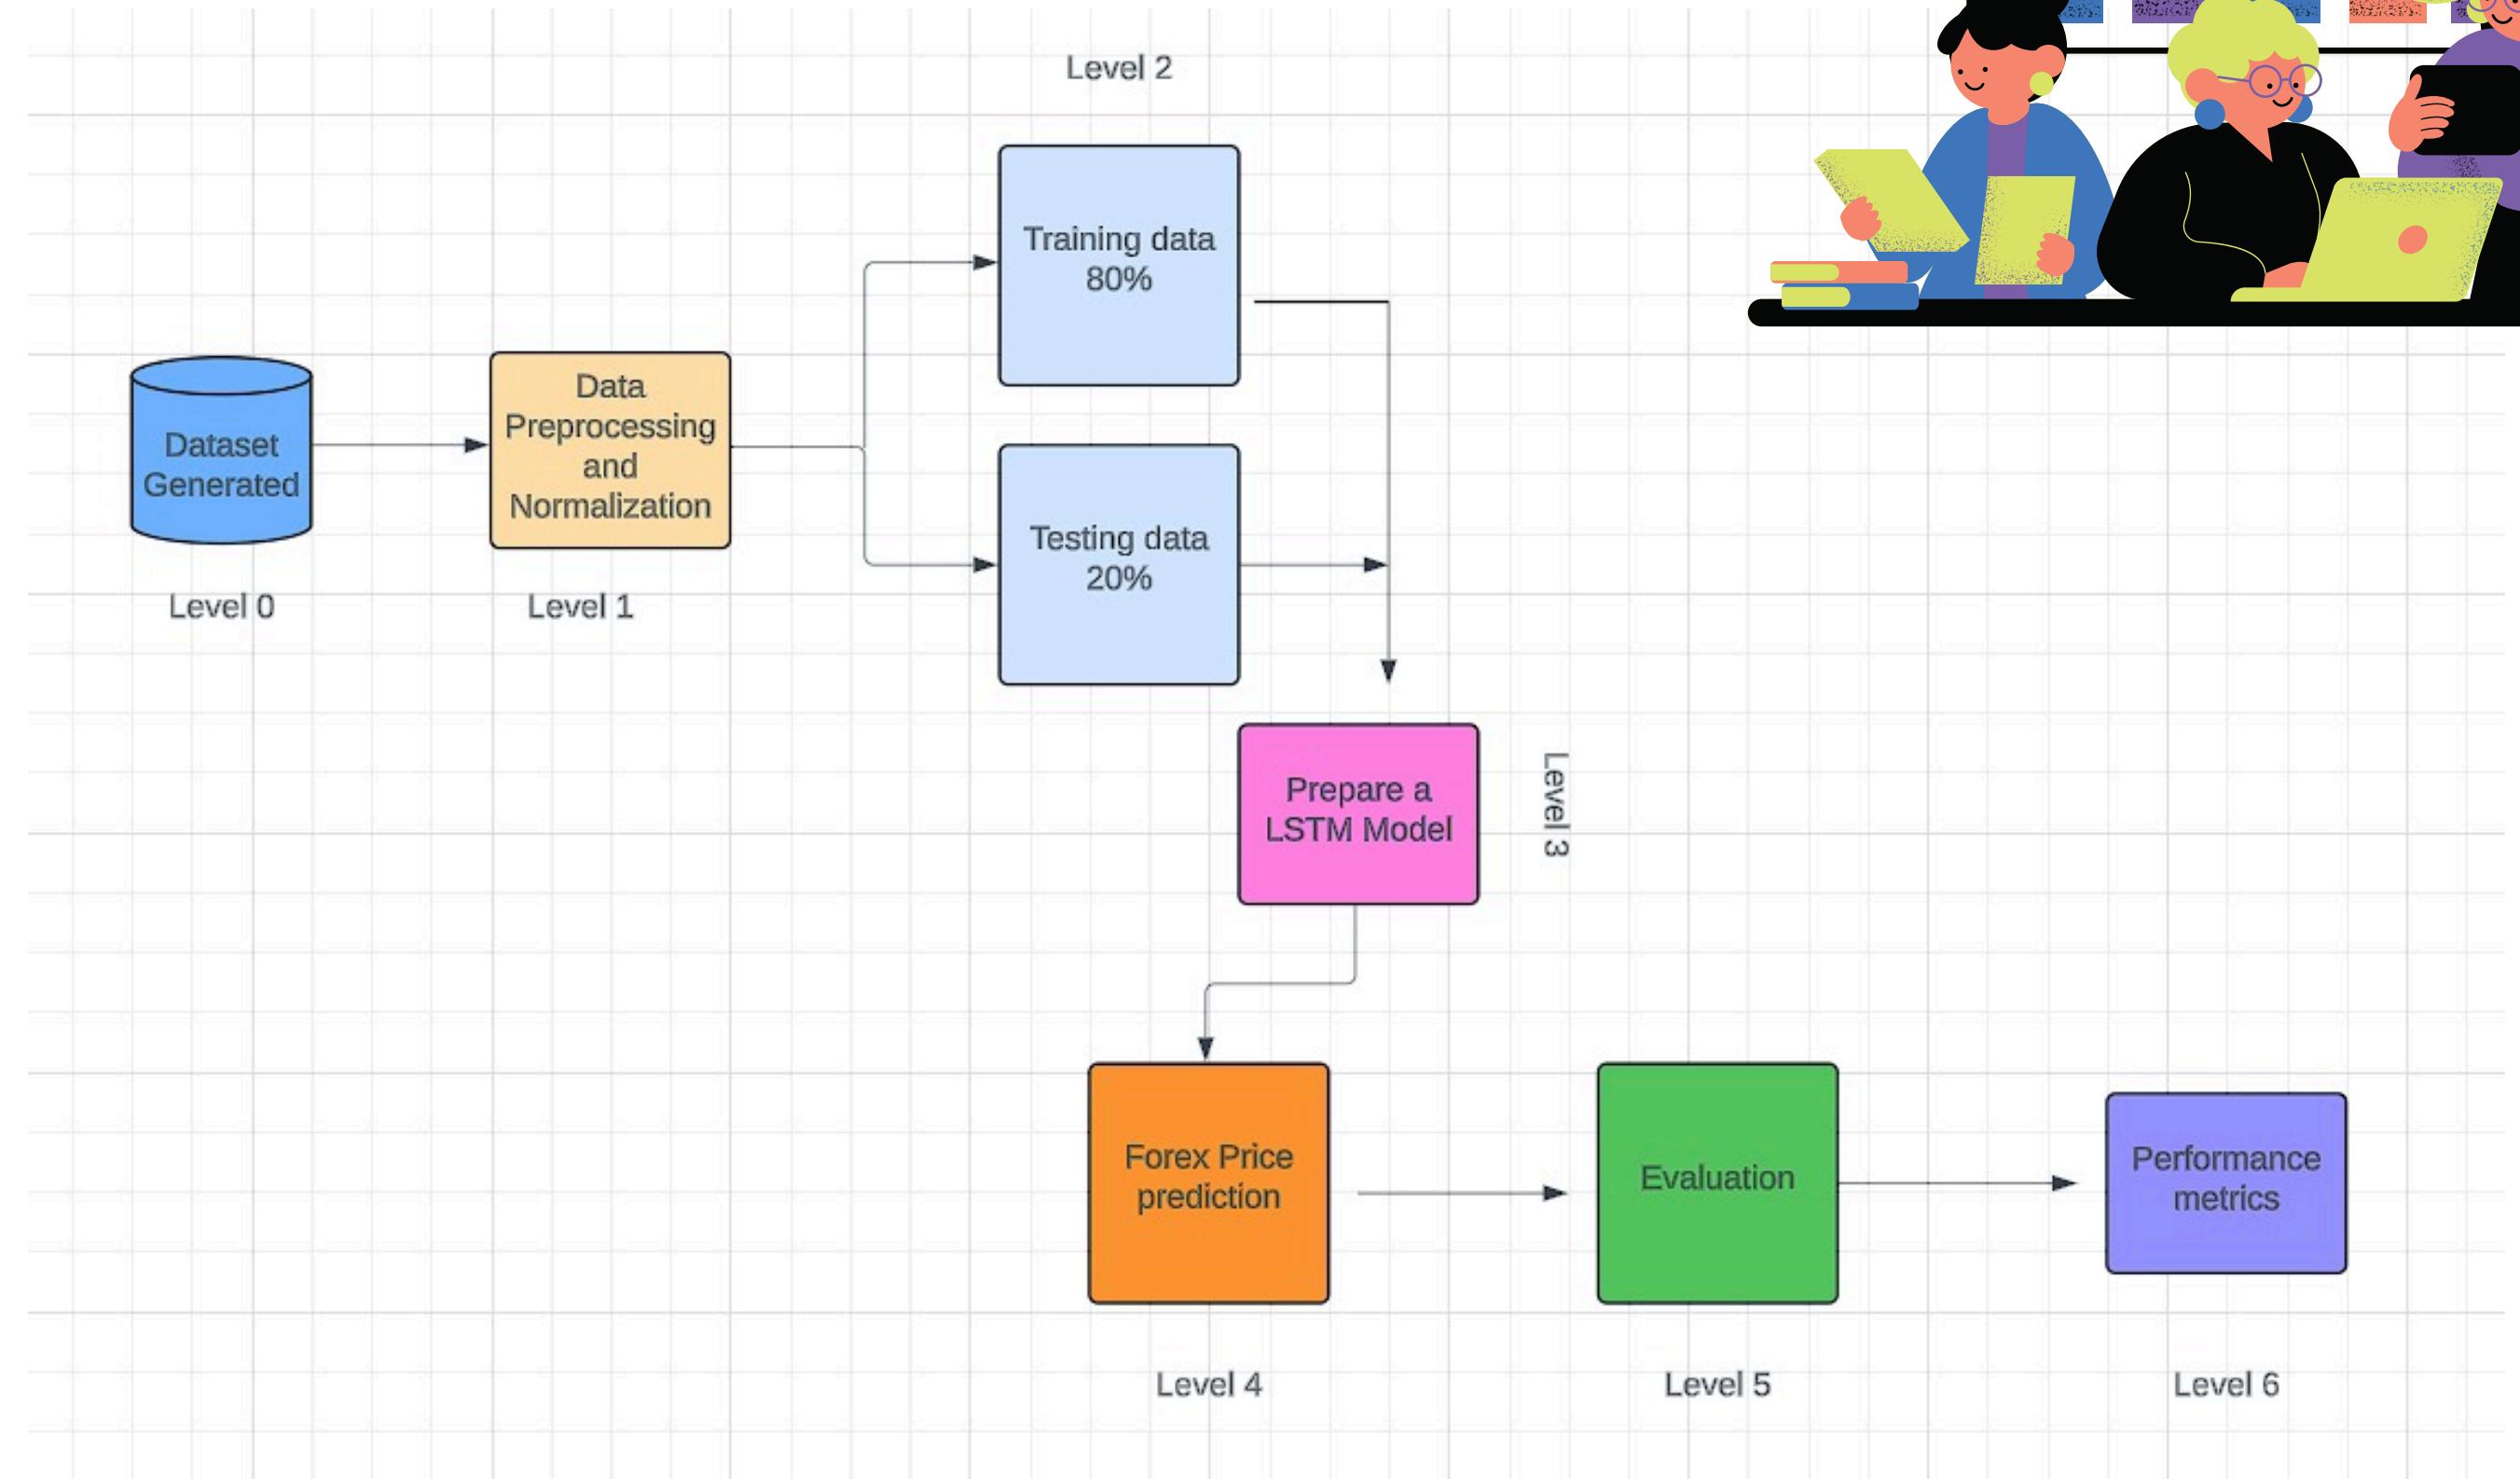


# ACTIVITY

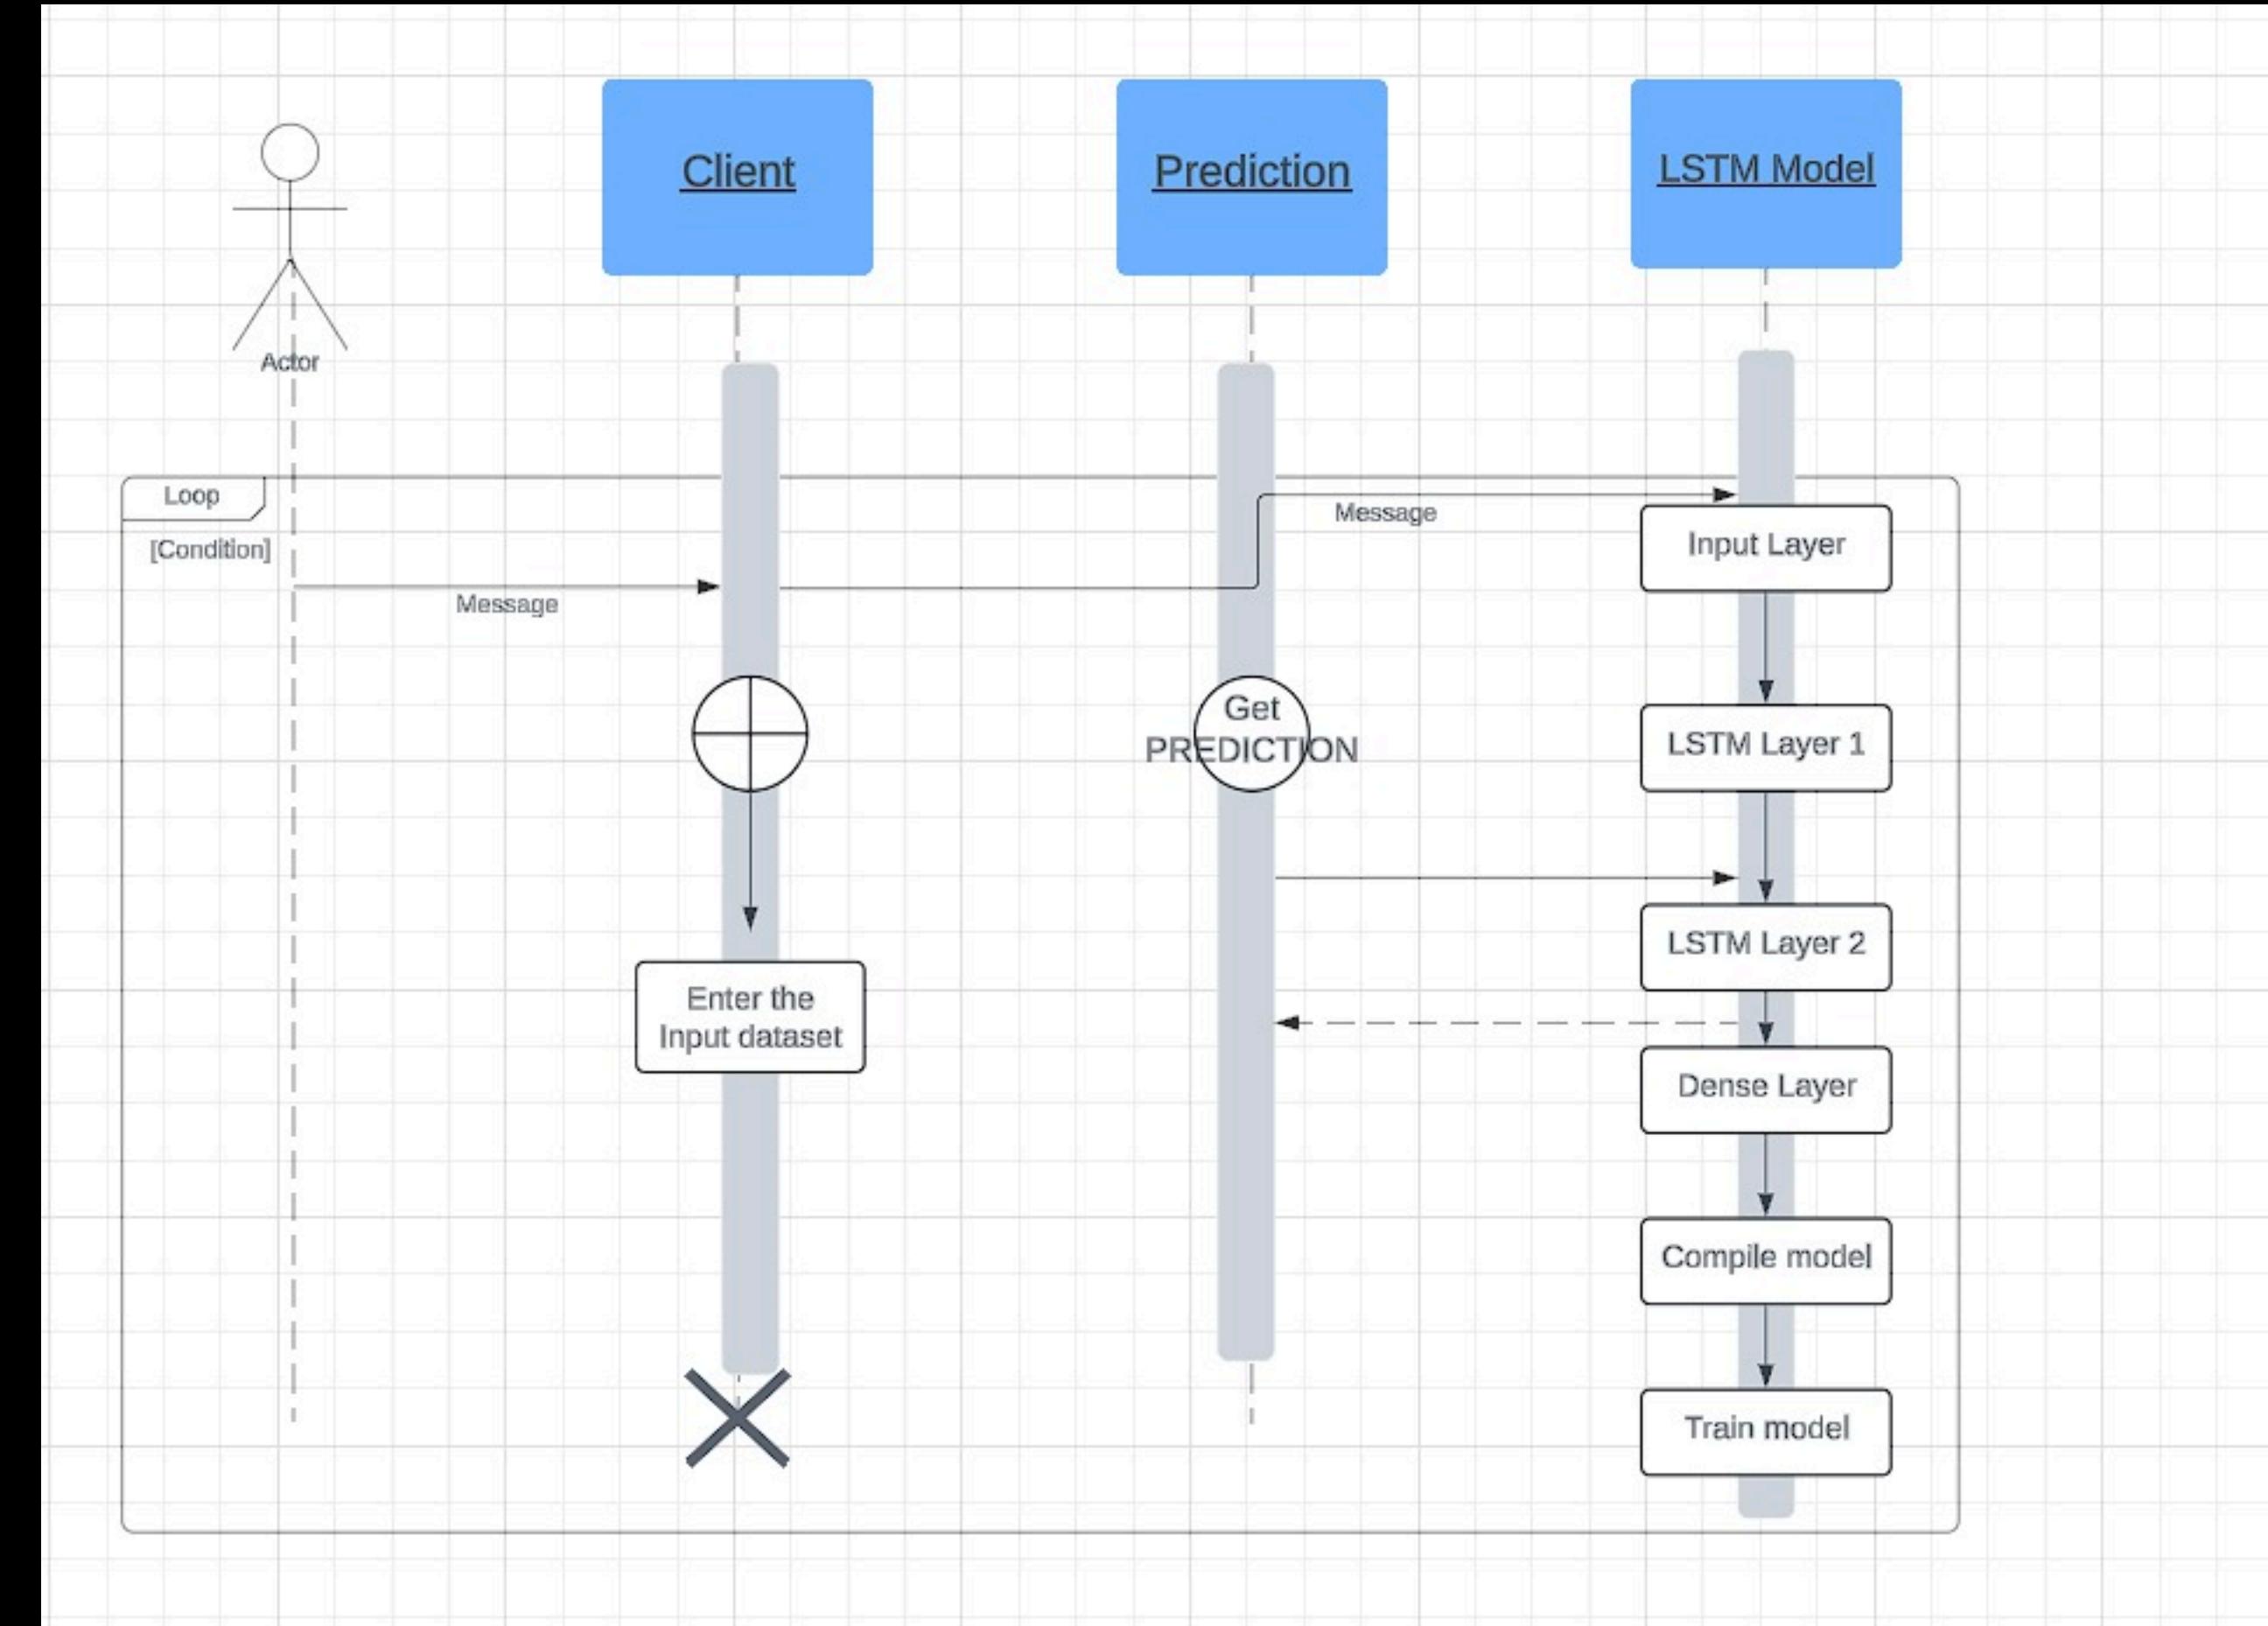
# CLASS DIAGRAM



# DATA FLOW DIAGRAM



# SEQUENCE DIAGRAM





# CODING

## Data pre-processing

Train-Test: 80:20

Synthetic and processed dataset

Samples: 5000

## Training and validation

100 epochs

NVIDIA GPU

## Model Architecture

LSTM

Sequence\_Length: 10

units: 50

learning rate: 0.0001

batch\_size: 32

## Evaluation Metrics

Various plots

loss



# IMPORT ALL NECESSARY LIBRARIES

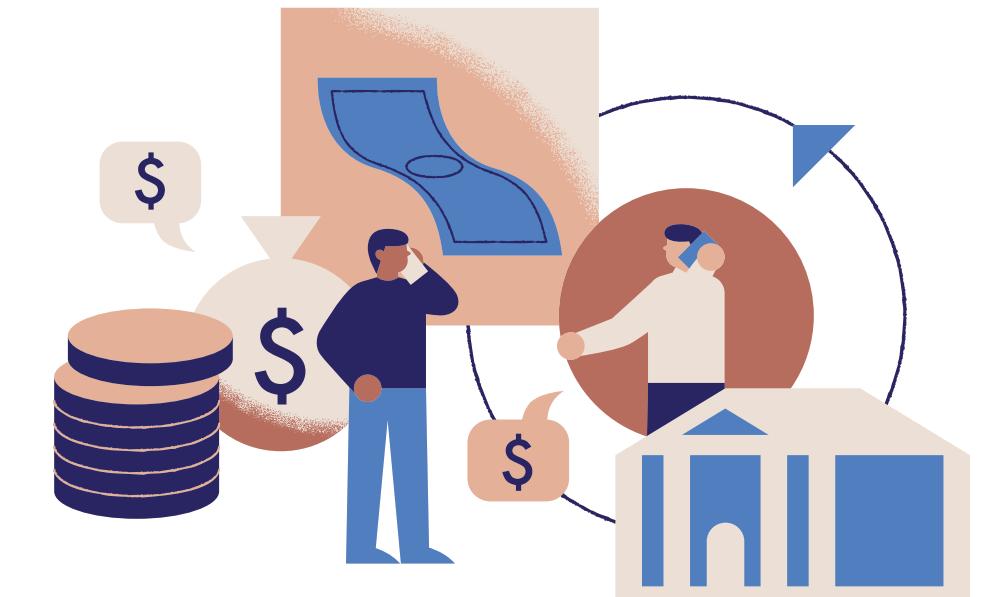
```
: import os  
import sys  
import json  
import time  
import pickle  
import numpy as np  
import pandas as pd  
from matplotlib import pyplot as plt
```

## DATASET

```
: train_data=pd.read_csv("C:\\\\Users\\\\91630\\\\Anaconda files\\\\New_Forex\\\\train_data.csv")  
test_data=pd.read_csv("C:\\\\Users\\\\91630\\\\Anaconda files\\\\New_Forex\\\\test_data.csv")
```

```
train_data.head()
```

Open	High	Low	Close	Volume	SMA	EMA	RSI	MACD	Upper_BB	Middle_BB	Lower_BB	%K	%D	ATR	Outcome
------	------	-----	-------	--------	-----	-----	-----	------	----------	-----------	----------	----	----	-----	---------



# HANDLING MISSING VALUES

```
import pandas as pd
from sklearn.preprocessing import MinMaxScaler

# ... (Load your data into train_data and test_data)

# Explicitly Exclude Specific Columns
columns_to_normalize = train_data.drop(columns=['Outcome', 'Upper_BB', 'Lower_BB', 'STD', 'Upper', 'Lower']).columns

scaler = MinMaxScaler()
scaler.fit(train_data[columns_to_normalize]) # Fit on train_data

# Transform Data
train_data_normalized = train_data.copy()
test_data_normalized = test_data.copy()

train_data_normalized[columns_to_normalize] = scaler.transform(train_data[columns_to_normalize])
test_data_normalized[columns_to_normalize] = scaler.transform(test_data[columns_to_normalize])

# Drop the first 19 rows with missing values from the training data
train_data_normalized = train_data_normalized.iloc[19:].copy()

print(train_data_normalized.isnull().sum())
print(test_data_normalized.isnull().sum())

Open      0
High      0
Low       0
Close     0
Volume    0
SMA       0
EMA       0
RSI       0
MACD     0
```

## TRAIN-TEST SPLIT 80:20

```
: print(train_data.shape)
print(test_data.shape)
```

```
(3984, 16)
(997, 16)
```

# CREATE SEQUENCES

```
: import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler
from keras.models import Sequential
from keras.layers import LSTM, Dense, Input
from keras.callbacks import EarlyStopping
from sklearn.metrics import mean_squared_error, r2_score
import tensorflow as tf
tf.debugging.enable_check_numerics()
from keras.optimizers import Adam

# ... (Load your data into train_data and test_data)

# 1. Data Preprocessing (Already Done - You Have Normalized Data)

# 2. Prepare Sequences for LSTM
def create_sequences(data, sequence_length):
    Xs, ys = [], []
    for i in range(len(data) - sequence_length):
        Xs.append(data.iloc[i:(i + sequence_length)].values)
        ys.append(data.iloc[i + sequence_length]['Outcome']) # Predict 'Outcome'
    return np.array(Xs), np.array(ys)

sequence_length = 10 # Adjust this based on your time series characteristics
X_train, y_train = create_sequences(train_data_normalized, sequence_length)
X_test, y_test = create_sequences(test_data_normalized, sequence_length)
```

# LET'S BUILD A LSTM MODEL



# MEAN IMPUTATION

```
# Check for NaNs in sequences
print(np.isnan(X_train).any())
print(np.isnan(X_test).any())

# Impute all NaNs in train_data_normalized and test_data_normalized
for col in train_data_normalized.columns:
    train_data_normalized[col] = train_data_normalized[col].fillna(train_data_normalized[col].mean())

for col in test_data_normalized.columns:
    test_data_normalized[col] = test_data_normalized[col].fillna(test_data_normalized[col].mean())

# ... (rest of your code)
```



```
: # Create sequences
X_train, y_train = create_sequences(train_data_normalized, sequence_length)
X_test, y_test = create_sequences(test_data_normalized, sequence_length)
# ...your code for checking and imputing missing values...
# Re-create sequences after imputation
X_train, y_train = create_sequences(train_data_normalized, sequence_length)
X_test, y_test = create_sequences(test_data_normalized, sequence_length)
```



# SEQUENTIAL MODEL TRAINING

```
n_features = train_data_normalized.shape[1] # Number of features

model = Sequential()
model.add(Input(shape=(sequence_length, n_features)))
model.add(LSTM(units=50, activation='relu', return_sequences=True))
model.add(LSTM(units=50, activation='relu'))
model.add(Dense(units=1, activation='sigmoid')) # Output layer for binary classification
model.compile(optimizer=Adam(learning_rate=0.0001, clipvalue=1.0), loss='binary_crossentropy', metrics=['accuracy'])
# 4. Train Model with Early Stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)
model.fit(X_train, y_train, epochs=100, batch_size=32, validation_split=0.2, callbacks=[early_stopping])
```

67%  
Profit 2021

## PREDICT AND EVALUATE

```
# Print training accuracy
train_accuracy = history.history['accuracy'][-1] # Get final training accuracy from the history object
print(f'Training Accuracy: {train_accuracy:.4f}')

# 5. Make Predictions and Evaluate (Modified)
y_pred_prob = model.predict(X_test) # Get probabilities
y_pred = (y_pred_prob > 0.5).astype(int) # Convert probabilities to binary labels (0 or 1)

# ... (inverse transform your predictions and labels if needed)

# Calculate and print test accuracy
test_accuracy = accuracy_score(y_test, y_pred)
print(f'Test Accuracy: {test_accuracy:.4f}')
```

```
Epoch 91/100  
99/99 4s 38ms/step - accuracy: 0.6755 - loss: 0.6038 - val_accuracy: 0.6497 - val_loss: 0.6310  
Epoch 92/100  
99/99 4s 39ms/step - accuracy: 0.6721 - loss: 0.6120 - val_accuracy: 0.6434 - val_loss: 0.6321  
Epoch 93/100  
99/99 4s 39ms/step - accuracy: 0.6554 - loss: 0.6090 - val_accuracy: 0.6497 - val_loss: 0.6403  
Epoch 94/100  
99/99 4s 38ms/step - accuracy: 0.6781 - loss: 0.6020 - val_accuracy: 0.6459 - val_loss: 0.6341  
Epoch 95/100  
99/99 4s 39ms/step - accuracy: 0.6719 - loss: 0.6051 - val_accuracy: 0.6371 - val_loss: 0.6493  
Epoch 96/100  
99/99 4s 38ms/step - accuracy: 0.6510 - loss: 0.6253 - val_accuracy: 0.6459 - val_loss: 0.6346  
Epoch 97/100  
99/99 4s 38ms/step - accuracy: 0.6676 - loss: 0.6035 - val_accuracy: 0.6472 - val_loss: 0.6353  
Epoch 98/100  
99/99 4s 40ms/step - accuracy: 0.6664 - loss: 0.6126 - val_accuracy: 0.6282 - val_loss: 0.6500  
Epoch 99/100  
99/99 4s 39ms/step - accuracy: 0.6655 - loss: 0.5986 - val_accuracy: 0.6561 - val_loss: 0.6379  
Epoch 100/100  
99/99 4s 38ms/step - accuracy: 0.6874 - loss: 0.5980 - val_accuracy: 0.6536 - val_loss: 0.6286
```



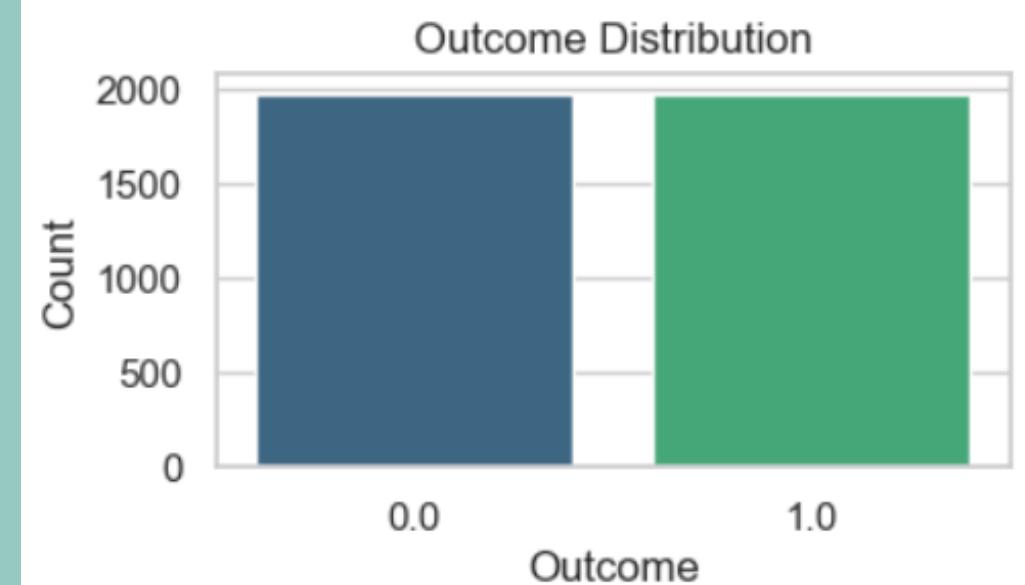
# EPOCHS

# VISUALIZATIONS

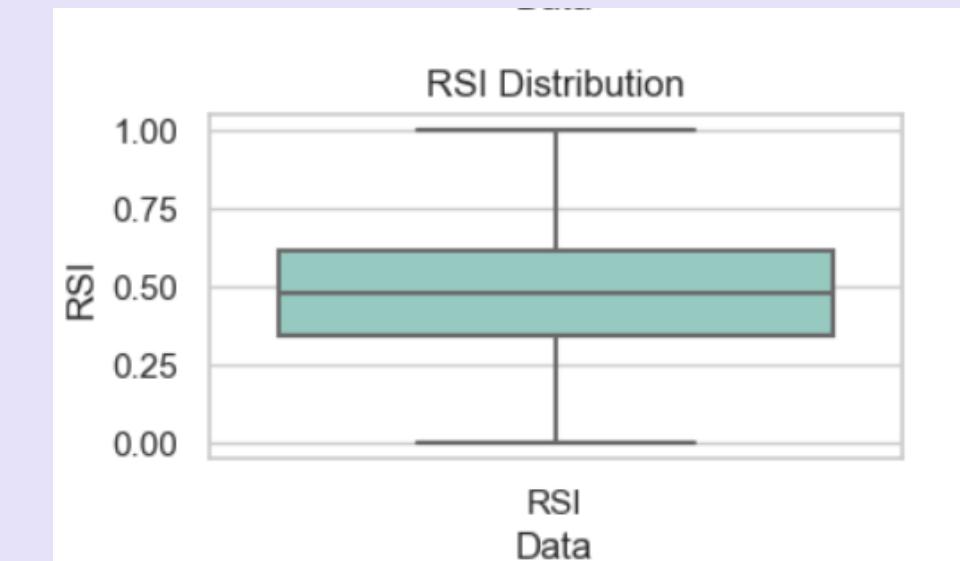
## HISTOGRAMS



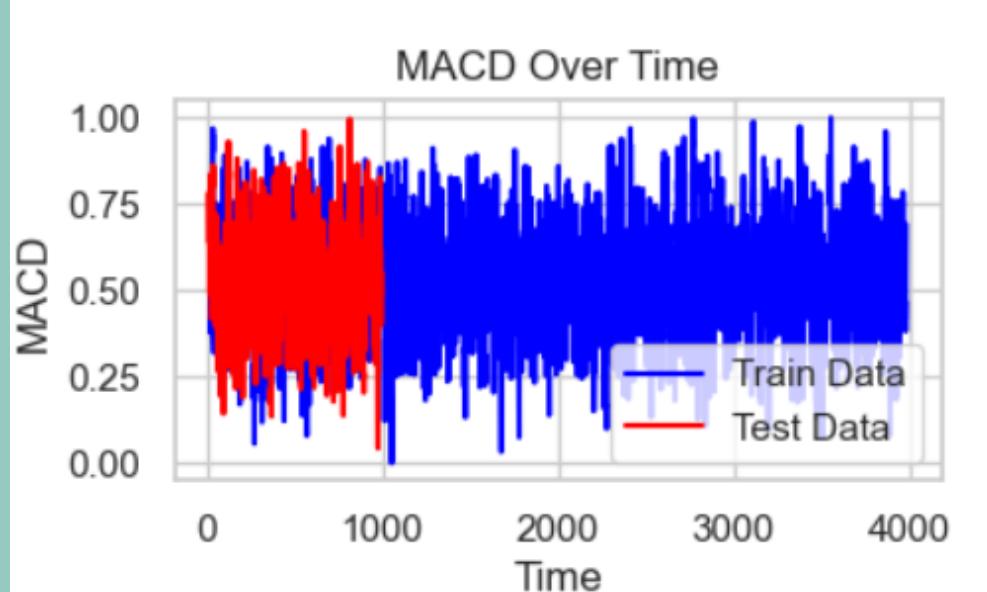
## BAR CHART



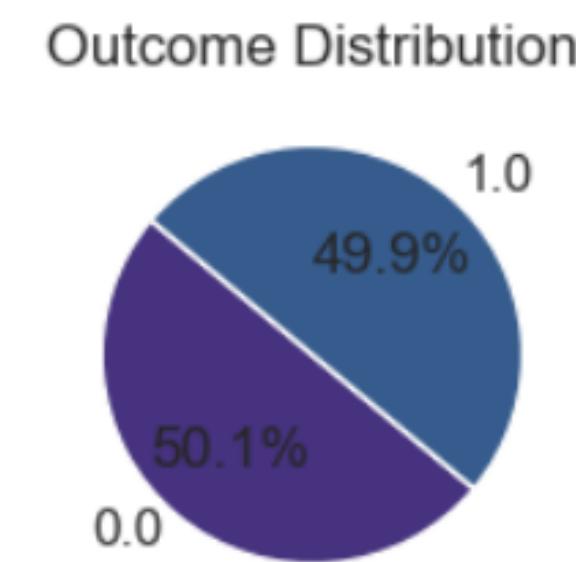
## BOX PLOT



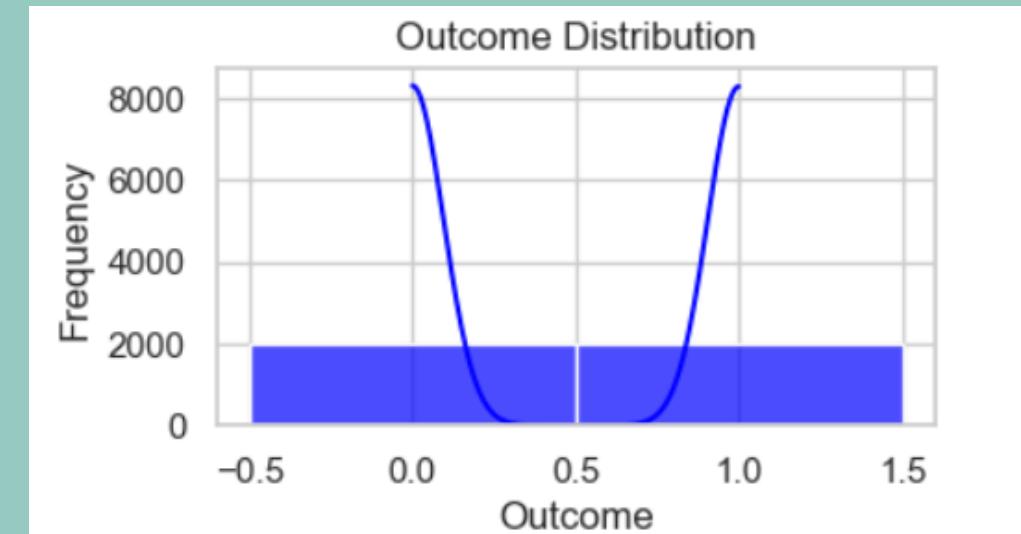
## LINE CHART



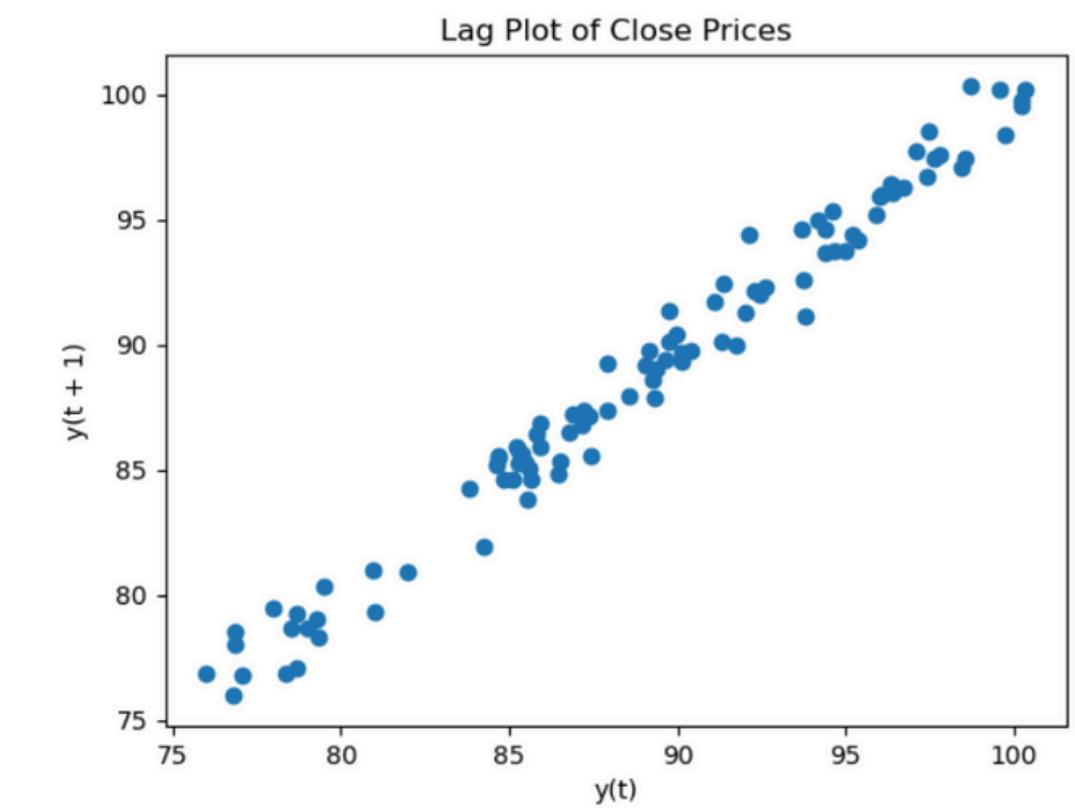
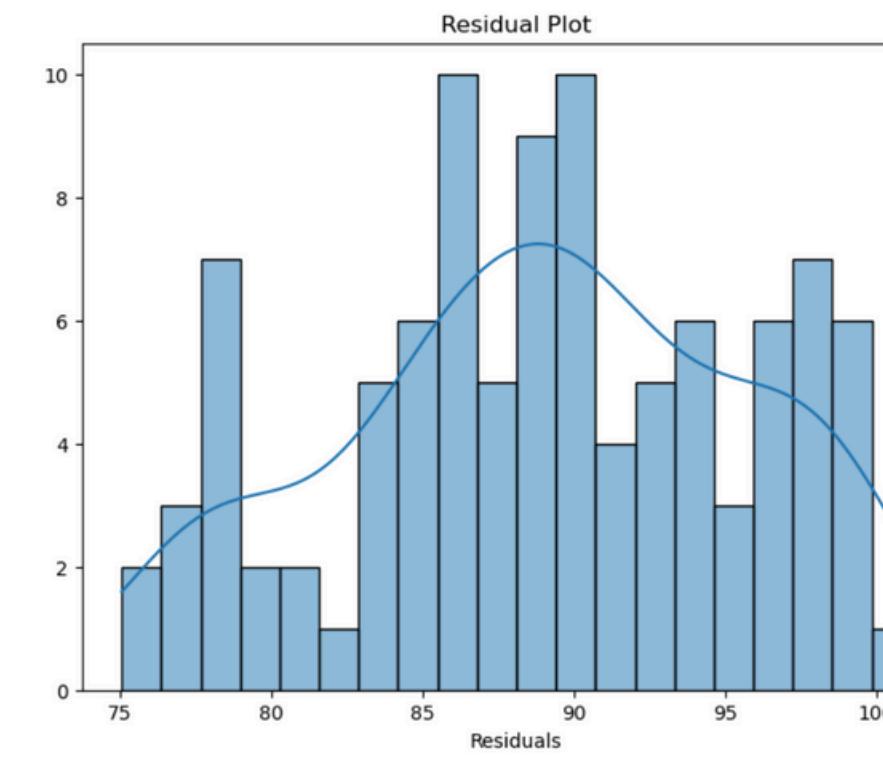
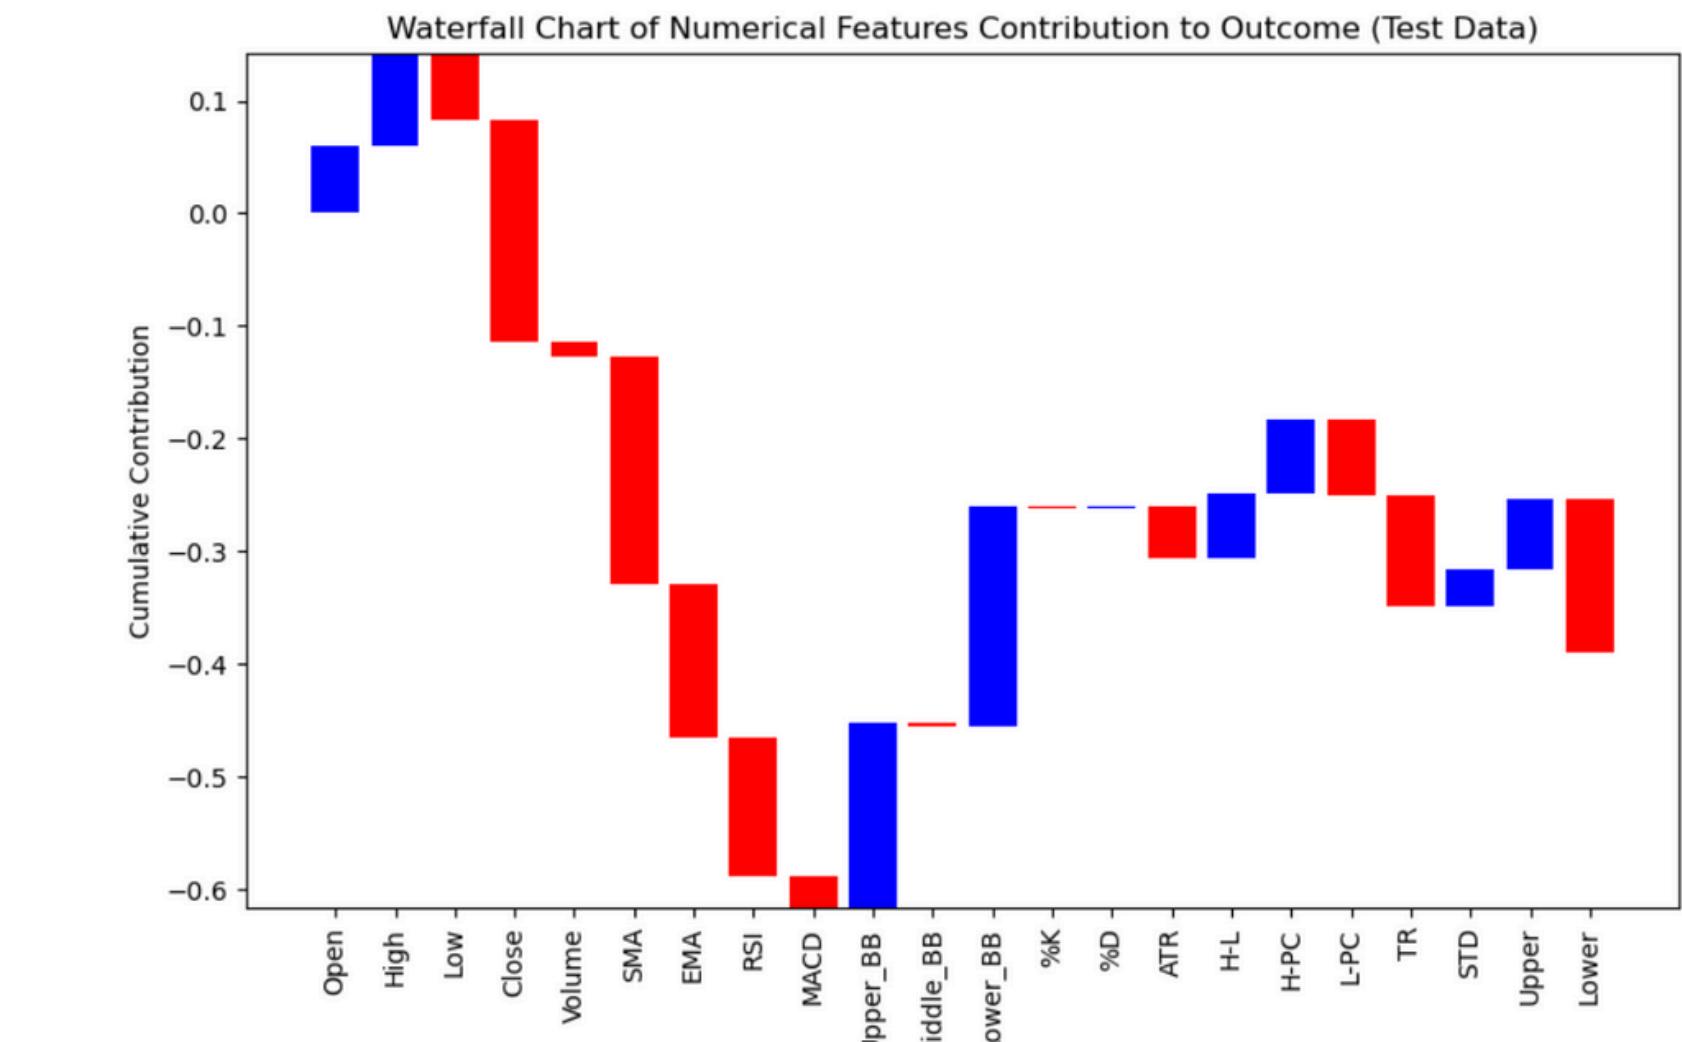
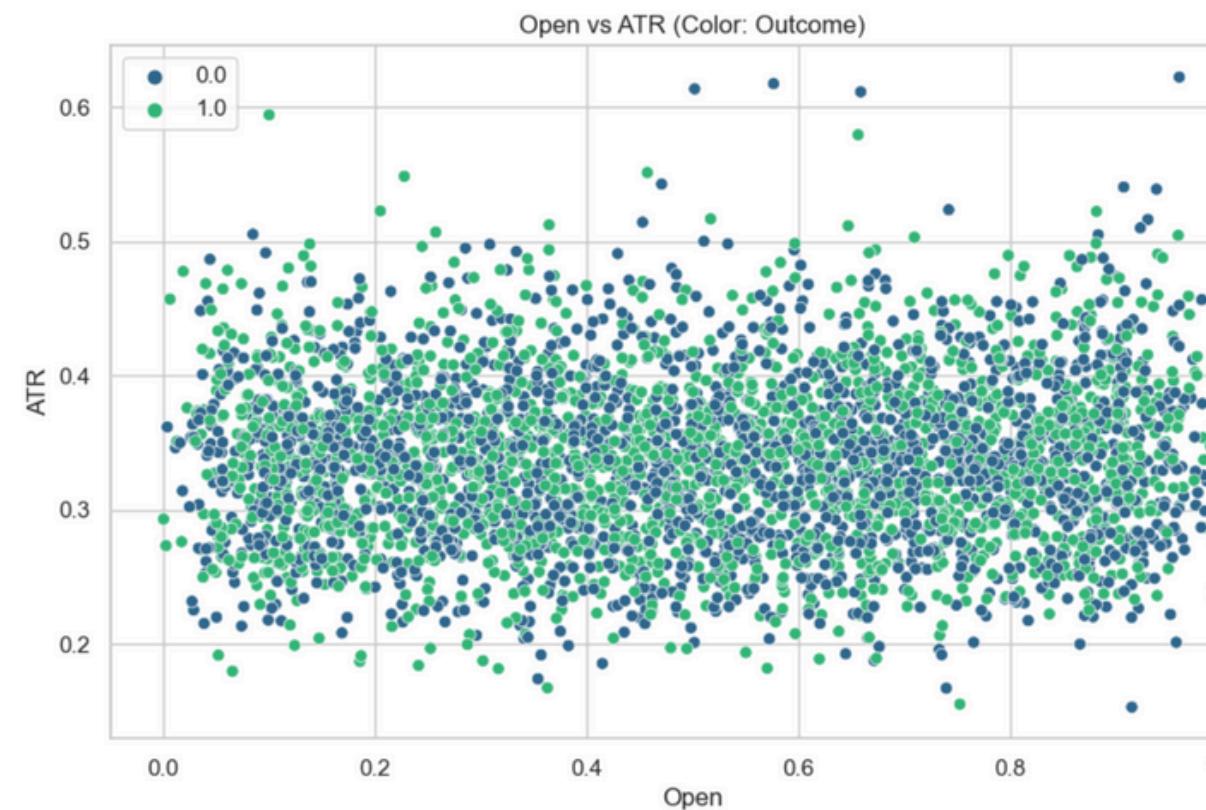
## PIE CHART



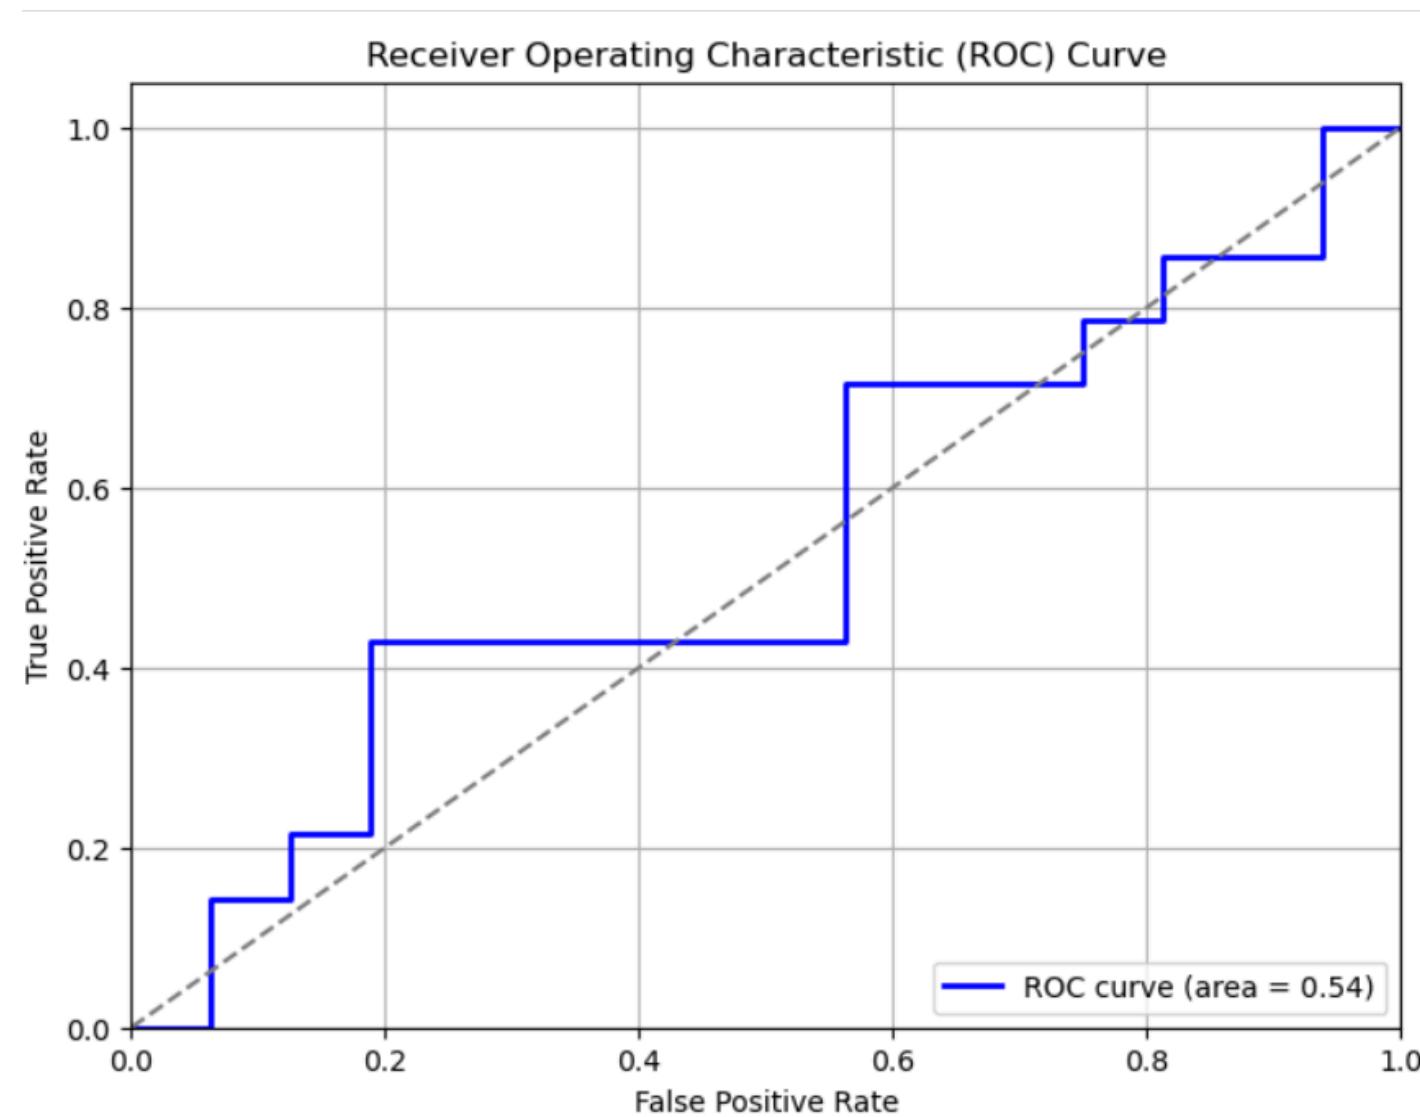
## AREA CHART



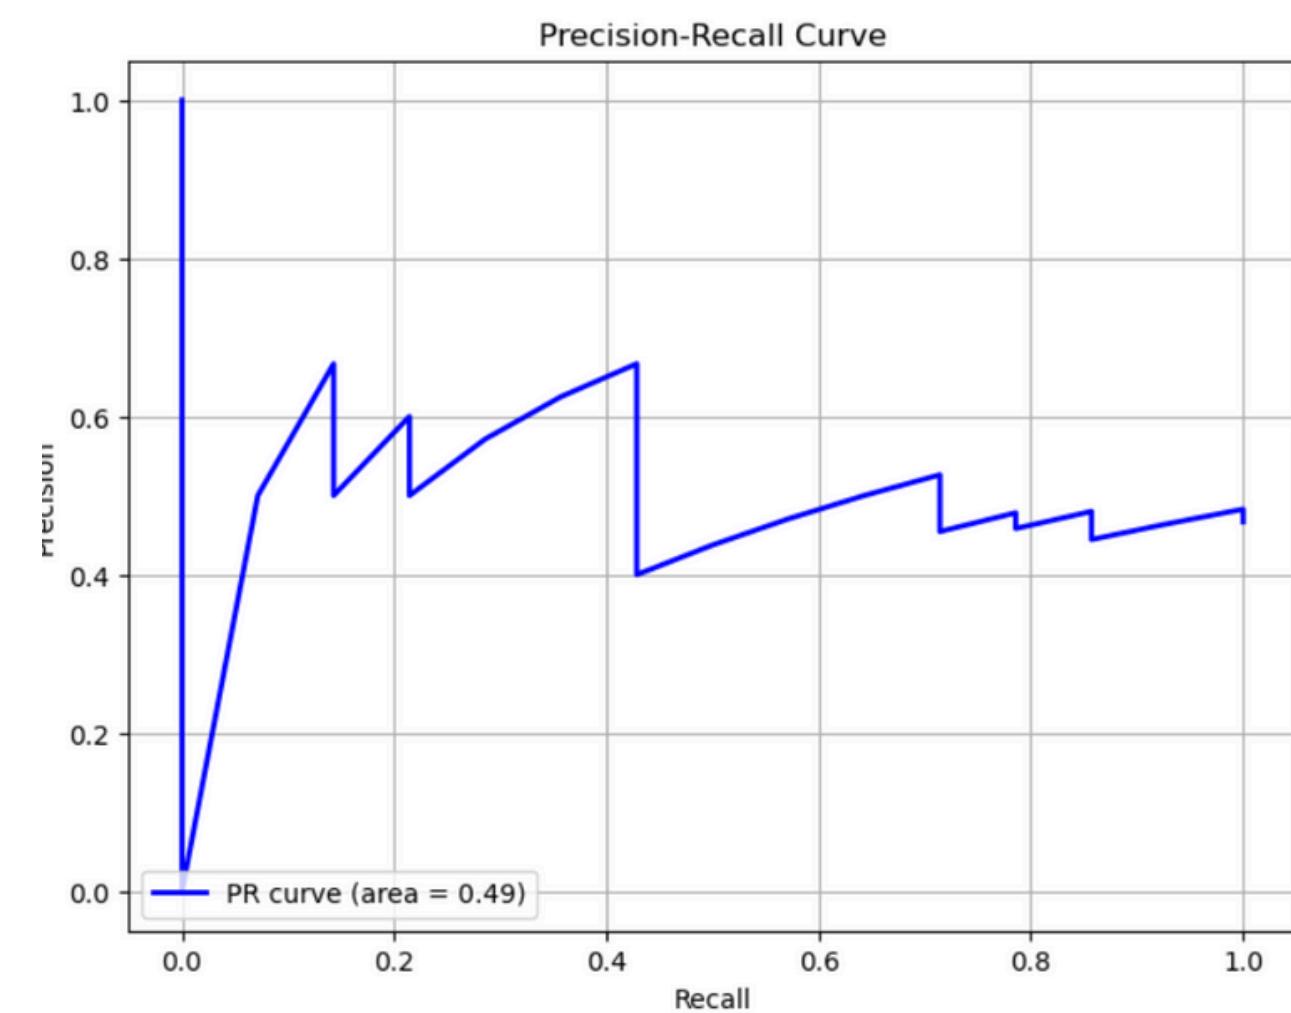
# PLOTTINGS



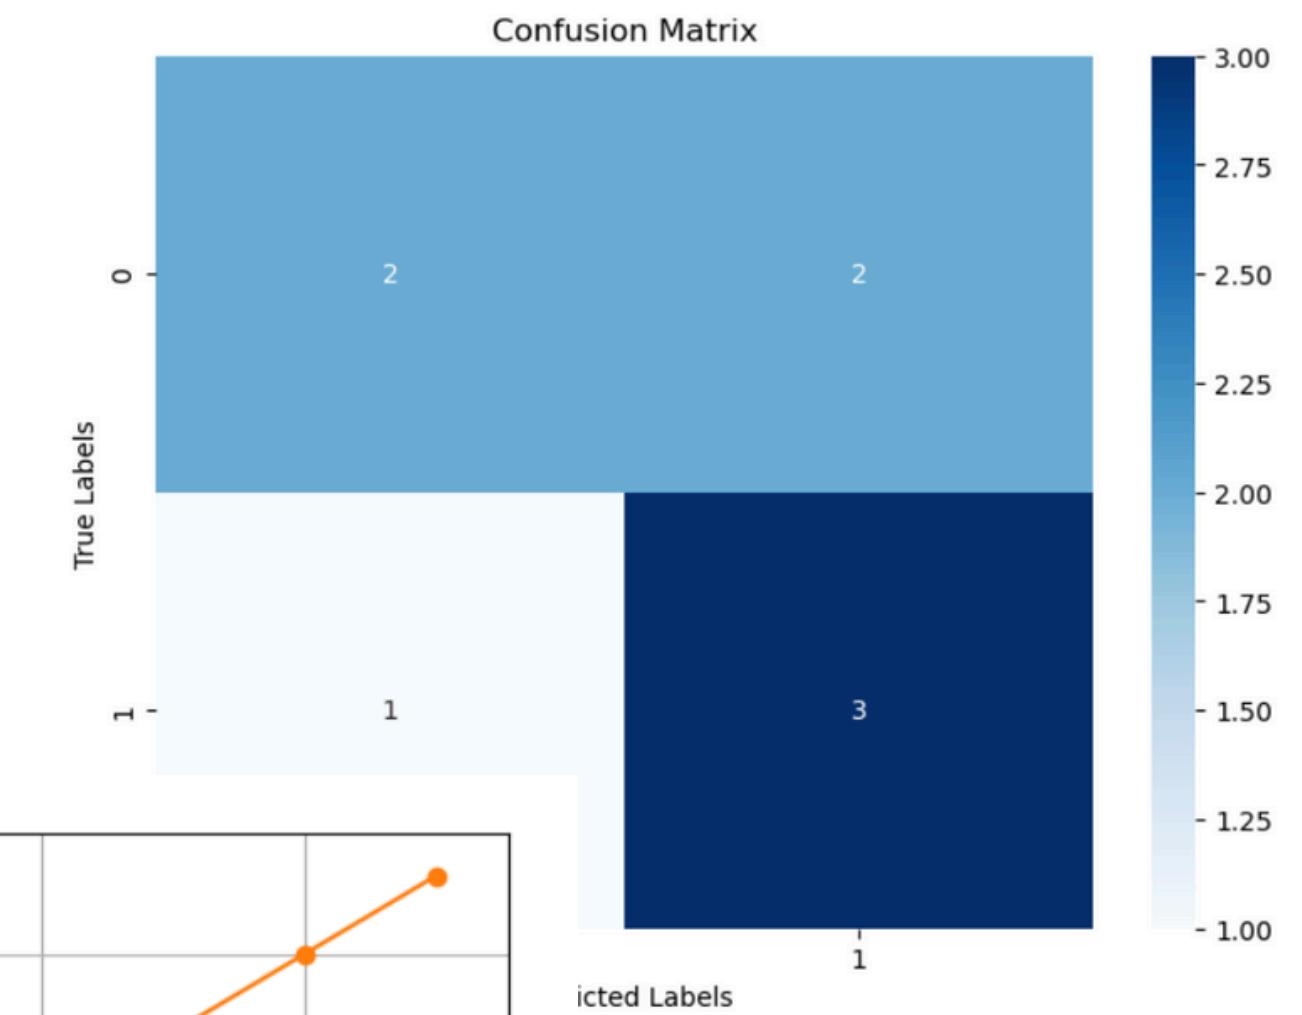
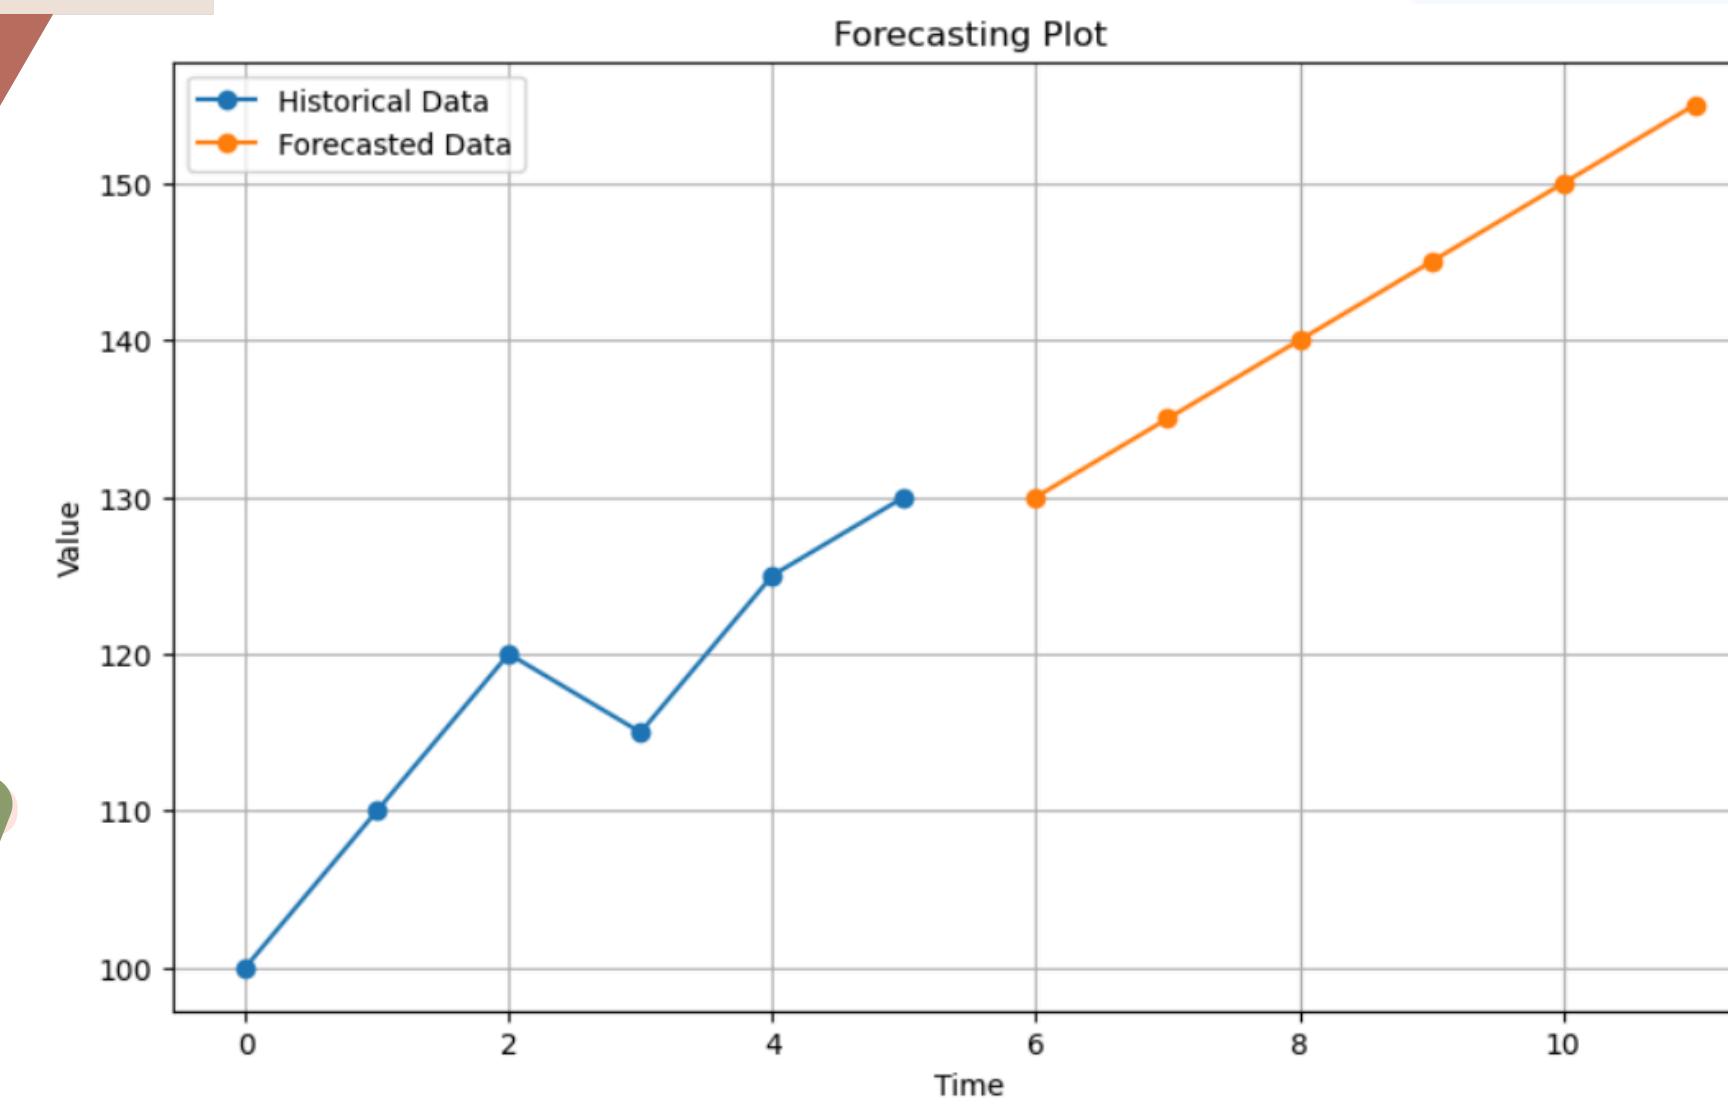
# ROC CURVE



# PR CURVE



# CONFUSION MATRIX



# TESTING: BLACKBOX

```
class LSTMBlackBoxTest(unittest.TestCase):
    def setUp(self):
        # Load and preprocess data
        self.X_train, self.X_test, self.y_train, self.y_test = load_and_preprocess_data()
        self.sequence_length = self.X_train.shape[1]
        self.input_shape = (self.X_train.shape[1], self.X_train.shape[2])
        self.model = create_lstm_model(self.input_shape)

    def test_valid_input(self):
        # Create a valid input
        valid_input = np.random.rand(1, self.sequence_length, self.input_shape[1])

        # Get the prediction
        prediction = self.model.predict(valid_input)

        # Assertions
        self.assertEqual(prediction.shape, (1, 1))
        self.assertTrue(0 <= prediction[0, 0] <= 1)

    # Function to run unittests
    def run_tests():
        suite = unittest.TestLoader().loadTestsFromTestCase(LSTMBlackBoxTest)
        unittest.TextTestRunner().run(suite)

    # Run the tests
    run_tests()
```

1/1 ————— 1s is/step

.

Ran 1 test in 1.796s

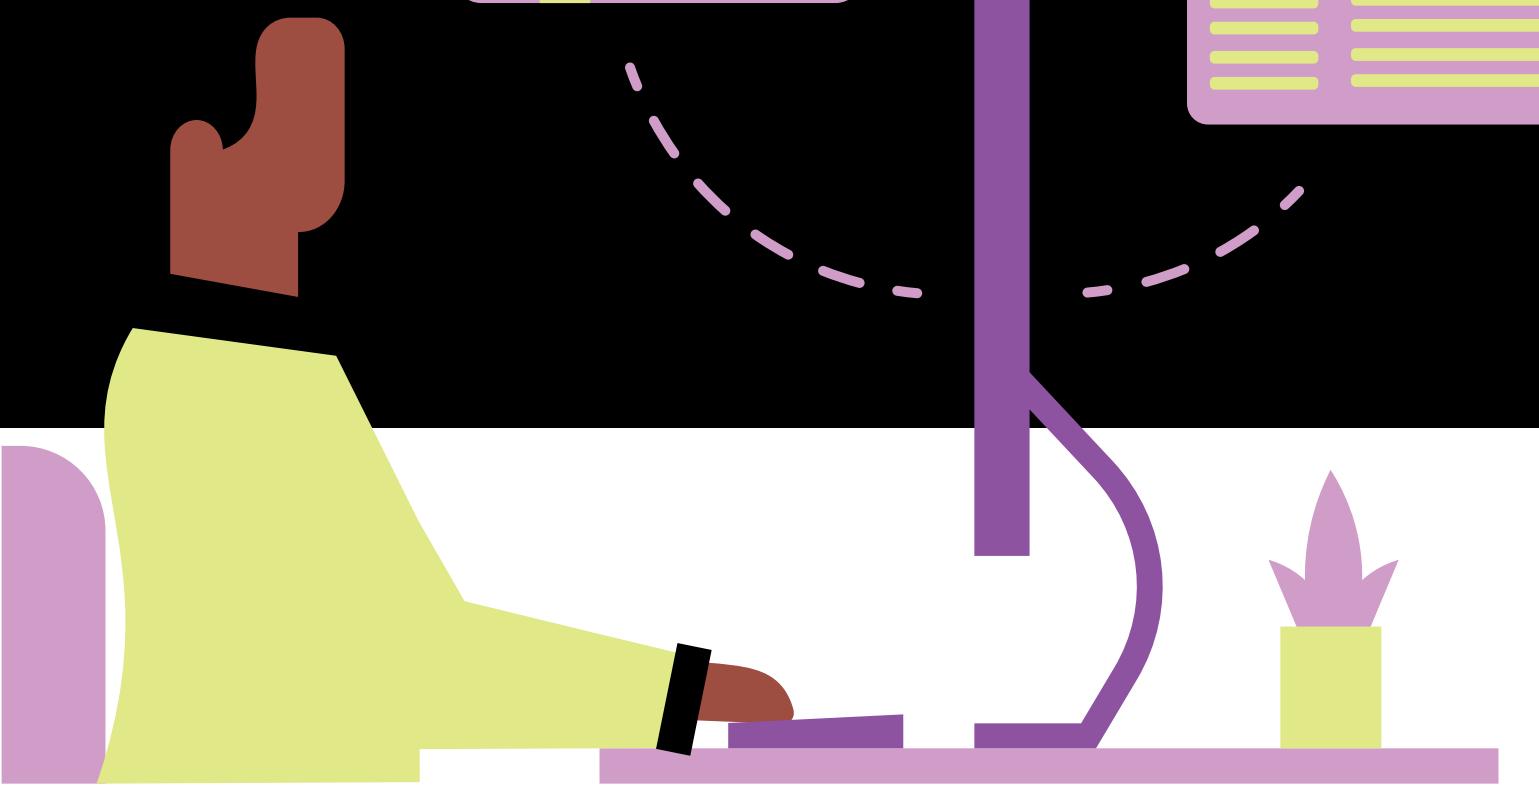
OK



# CONCLUSION



Continued advancements in LSTM-based forecasting methodologies hold promise for further improving decision-making processes in financial markets, albeit with ongoing challenges in real-world implementation and deployment. These conclusions underscore LSTM's role as a powerful tool in forecasting directional movements, with ongoing research aimed at refining its capabilities and addressing current limitations.



THANK YOU